Explaining Predictions from Machine Learning Models: Algorithms, Humans, and Pedagogy

Ana Lučić

Explaining Predictions from Machine Learning Models: Algorithms, Humans, and Pedagogy

ACADEMISCH PROEFSCHRIFT

ter verkrijging van de graad van doctor aan de
Universiteit van Amsterdam
op gezag van de Rector Magnificus
prof. dr. ir. K.I.J. Maex
ten overstaan van een door het College voor Promoties ingestelde
commissie, in het openbaar te verdedigen in
de Aula der Universiteit
op 23 dag 10 maand 2012, te 14:00 uur

door

Ana Lučić

geboren te Paraćin, Joegoslavië

Promotiecommissie

Promotor:

Prof. dr. M. de Rijke Universiteit van Amsterdam

Co-promotor:

Prof. dr. H. Haned Universiteit van Amsterdam

Overige leden:

Prof. dr. M. Lovrić McMaster University

Prof. dr. C. Sánchez Gutiérrez
Dr. F. P. Santos
Universiteit van Amsterdam
Universiteit van Amsterdam
Sapienza University of Rome
Prof. dr. M. Welling
Universiteit van Amsterdam
Universiteit van Amsterdam

Faculteit der Natuurwetenschappen, Wiskunde en Informatica

This research was supported by Ahold Delhaize and the Netherlands Organisation for Scientific Research under project number 652.001.003.

Copyright © 2022 Ana Lučić, Amsterdam, The Netherlands Cover by Someone Else Printed by Printer Name, City

ISBN: xxx-xx-xxxx-xxx-x

The more I see, the less I know.

– Anthony Kiedis

Contents

1	Intr	duction	1
	1.1	Research Outline and Questions	1
	1.2	Main Contributions	2
	1.3	Thesis Overview	2
	1.4	Origins	3
I	Alg	prithms	5
2	Cou	terfactual Explanations for Tree Ensembles	7
	2.1		7
	2.2		9
		2.2.1 Counterfactual Explanations	9
		2.2.2 Algorithmic Recourse	9
		2.2.3 Adversarial Examples	9
		2.2.4 Differentiable Tree-based Models)
	2.3	Problem Formulation)
	2.4	Method: FOCUS	1
		2.4.1 Loss Function Definitions	1
		2.4.2 Tree-based Models	2
		2.4.3 Approximations of Tree-based Models	3
		2.4.4 Our Method: FOCUS	4
		2.4.5 Effects of Hyperparameters	4
	2.5	Experimental Setup	4
		2.5.1 Datasets	4
		2.5.2 Models	5
		2.5.3 Distance Functions	5
		2.5.4 Evaluation Metrics	5
	2.6	Experiment 1: FOCUS vs. FT	5
		2.6.1 Baseline: Feature Tweaking	5
		2.6.2 Results	7
	2.7	Experiment 2: FOCUS vs. DACE	9
		2.7.1 Baseline: DACE	9
		2.7.2 Results	0
	2.8	Discussion and Analysis	O
		2.8.1 Case Study: Credit Risk	1
	2.9	Conclusion	4
3	Con	terfactual Explanations for Graph Neural Networks 22	7
-	3.1	Introduction	
	3.2	Related Work	
		3.2.1 GNN Explainability	
		3.2.2 Counterfactual Explanations	
		3.2.3 Adversarial Attacks	
		3	-

CONTENTS

	3.3	Backgı	round	30
		3.3.1	Graph Neural Networks	30
		3.3.2	Matrix Sparsification	31
	3.4	Proble	m Formulation	31
	3.5	Metho	d: CF-GNNExplainer	32
		3.5.1	Adjacency Matrix Perturbation	32
		3.5.2	Counterfactual Generating Model	32
		3.5.3	Loss Function Optimization	33
		3.5.4	CF-GNNExplainer	33
		3.5.5	Complexity	34
	3.6	Experi	mental Setup	34
		3.6.1	Datasets and Models	35
		3.6.2	Baselines	36
		3.6.3	Metrics	36
		3.6.4	Hyperparameter Search	37
	3.7	Results	S	37
		3.7.1	Main Findings	37
		3.7.2	Comparison to GNNExplainer	38
		3.7.3	Summary of results	38
	3.8	Societa	al Impact	40
	3.9	Conclu	ision	41
II	Hı	ımans	3	43
	0			4 -
4			1	45
	4.1		uction	45
	4.2		d Work	48
	4.3		d	50
	4.4	-	mental Setup	52
		4.4.1	Dataset and Model	52
		4.4.2	Comparison to LIME	53
		4.4.3	User Study Design	54
	4.5		mental Results	56
		4.5.1	Objective Questions	56
		4.5.2	Subjective Questions	57
	4.6		sion	57
		4.6.1	Comparing Attitudes Conditioned on Background	58
		4.6.2	Comparing Attitudes in the Treatment Group	59
		4.6.3	User Study Limitations	59
	4.7	Conclu	ısion	60

II	I P	edago	gy	63
5	Teac	ching R	Responsible AI through Reproducibility	65
	5.1	_	uction	. 65
	5.2	Relate	d Work	. 67
		5.2.1	Characterizing Responsible AI Courses	. 67
		5.2.2	Similar Responsible AI Courses	
	5.3	Integra	ating Reproducibility of AI Research into the FACT-AI Course	
	5.4	Learni	ing Objectives	. 69
	5.5	Course	e Setup	. 70
		5.5.1	Lectures	. 70
		5.5.2	Paper Discussions	. 71
		5.5.3	Group Project	. 71
		5.5.4	Taking the Course Online	. 73
	5.6	Feedba	ack	. 73
		5.6.1	Feedback from Students	. 73
		5.6.2	Feedback from the ML Reproducibility Challenge	. 74
	5.7	Factor	s Contributing to a Successful Course	. 75
		5.7.1	General	. 75
		5.7.2	Concerning FACT-AI	. 77
		5.7.3	Concerning Reproducibility	. 77
	5.8	Areas	of Improvement	
		5.8.1	General	
		5.8.2	Concerning FACT-AI	. 78
		5.8.3	Concerning Reproducibility	
	5.9	Concl	usion	. 78
6	Con	clusior	18	81
	6.1	What	Have We Done?	. 81
	6.2	What	Is Next?	. 81
Bi	bliog	raphy		83
Su	mma	ıry		89
Sa	žetak			91
Sa	menv	atting		93

Introduction

- What is machine learning?
- How can we do ML responsibly? What are aspects of responsible ML (F-A-C-T).
- Relate transparency to explainability. What is explainability and why is it important?
- What kinds of explanations should we be trying to generate? Explanations that are contrastive/counterfactual are intuitive for humans [Miller, 2019] (to flow into first two RQs about counterfactual explanations, then RQ3 about humans). What is a contrastive/counterfactual explanation?
- Summary of what we do in the thesis: introduce a method for generating counterfactual explanations for tree ensembles, then extend this to graph-based models, then look at real use case + real users, and design a method for explaining errors in regression predictions which is evaluated with a user study, then transition from creating knowledge to translating knowledge in the FACT-AI course paper.

1.1 RESEARCH OUTLINE AND QUESTIONS

This thesis focuses on explaining predictions for ML models from three different contexts/positions: (i) algorithms, (ii) humans, (iii) pedagogy.

- **RQ1** How can we generate counterfactual explanations specifically for tree-based models?
- RQ2 How can we extend our explanation method for tree-based models to graph-based models?
- **RQ3** How can we develop an explanation method based on a real-world use case and evaluate it in a human-centric way?

RQ4 How can we teach about responsible AI topics to a technical, research-oriented audience?

1.2 MAIN CONTRIBUTIONS

In this section, we summarize the main contributions of this thesis.

Theoretical Contributions

- 1. A formalization of the counterfactual explanation problem for GNNs (Chapter 3).
- 2. An experimental setup for evaluating counterfactual explanations for GNNs (Chapter 3).
- 3. A user study framework for evaluating the effectiveness of contrastive explanations (Chapter 4).
- 4. An analysis on the difference in attitudes towards explanations between different types of stakeholders (Chapter 4).

Algorithmic Contributions

- 5. Flexible Optimizable CoUnterfactual Explanations for Tree EnsembleS (FOCUS): an algorithm for generating counterfactual explanations for tree ensembles (Chapter 2).
- 6. CF-GNNExplainer: an algorithm for generating counterfactual explanations for graph neural networks (Chapter 3).
- 7. Monte Carlo Bounds for Reasonable Predictions (MC-BRP): an algorithm for generating explanations about errors in regression predictions (Chapter 4).

Pedagogical Contributions

8. A teaching setup for a course about responsible AI, with a focus on reproducibility (Chapter 5).

1.3 THESIS OVERVIEW

This thesis is organized into three parts.

The first part focuses on proposing new algorithms for explaining predictions from ML models. Specifically, we propose methods for generating counterfactual explanations for tree-based models (Chapter 2), and graph-based models (Chapter 3). These methods can be applied on any tree- or graph-based model, respectively.

The second part focuses on the interaction between explanations and humans. We propose a method for explaining errors in regression predictions based on a real use case at a large Dutch retailer (Chapter 4). We evaluate our method through a user

study on both practitioners from the retailer as well as researchers from the University of Amsterdam.

In the third part, we shift our focus from creating knowledge to translating knowledge to the next generation of researchers. We propose a course setup for teaching about responsible AI topics to a master-level audience and reflect on our learnings from past implementations of the course at the University of Amsterdam (Chapter 5).

1.4 ORIGINS

Below we list the publications that are the origins of each chapter.

Chapter 3 is based on the following paper:

 Ana Lucic, Harrie Oosterhuis, Hinda Haned, and Maarten de Rijke. FOCUS: Flexible Optimizable Counterfactual Explanations for Tree Ensembles. In AAAI Conference on Artificial Intelligence, 2022b.

AL and HO designed the method. AL ran the experiments. All authors contributed to the writing, AL did most of the writing.

Chapter 4 is based on the following paper:

 Ana Lucic, Maartje ter Hoeve, Gabriele Tolomei, Maarten de Rijke, and Fabrizio Silvestri. CF-GNNExplainer: Counterfactual Explanations for Graph Neural Networks. In *International Conference on Artificial Intelligence* and Statistics, 2022c.

AL designed the method and ran the experiments. All authors contributed to the writing, AL did most of the writing.

Chapter 5 is based on the following paper:

Ana Lucic, Hinda Haned, and Maarten de Rijke. Why Does My Model Fail?
 Contrastive Local Explanations for Retail Forecasting. In ACM Conference on Fairness, Accountability, and Transparency, 2020.

AL designed the method and ran the experiments. All authors contributed to the writing, AL did most of the writing.

Chapter 6 is based on the following paper:

 Ana Lucic, Maurits Bleeker, Sami Jullien, Samarth Bhargav, and Maarten de Rijke. Reproducibility as a Mechanism for Teaching Fairness, Accountability, Confidentiality, and Transparency in Artificial Intelligence. In AAAI Symposium on Educational Advances in Artificial Intelligence, 2022a.

AL and MB designed the course and implemented it together with MdR. All authors contributed to the writing, AL did most of the writing.

The writing of the thesis also benefited from work on the following publications:

- Ana Lucic, Hinda Haned, and Maarten de Rijke. Explaining Predictions from Tree-based Boosting Ensembles. In SIGIR Workshop on Fairness, Accountability, Confidentiality, and Transparency in Information Retrieval, 2019.
- Kim de Bie, Ana Lucic, and Hinda Haned. To Trust or Not to Trust a Regressor: Estimating and Explaining Trustworthiness of Regression Predictions. In *ICML Workshop on Human in the Loop Learning*, 2021.
- Michael Neely, Stefan F. Schouten, Maurits J. R. Bleeker, and Ana Lucic. Order in the Court: Explainable AI Methods Prone to Disagreement. In *ICML Workshop* on Theoretic Foundation, Criticism, and Application Trend of Explainable AI, 2021.
- Ana Lucic, Madhulika Srikumar, Umang Bhatt, Alice Xiang, Ankur Taly, Q. Vera Liao, and Maarten de Rijke. A Multistakeholder Approach Towards Evaluating AI Transparency Mechanisms. In CHI Workshop on Operationalizing Human-Centered Perspectives in Explainable AI, 2021.
- Surya Karunagaran, Ana Lucic, and Christine Custis. XAI Toolsheets: An Evaluation Framework for XAI Tools. In *Under review at AAAI/ACM Conference on AI, Ethics, and Society,* 2022. **Update when/if it's accepted**
- Add PAI paper once submitted.

Part I Algorithms

2

Counterfactual Explanations for Tree Ensembles

In this part of the thesis, we explore creating algorithms for explaining predictions from various types of ML models. In this chapter, we address the following research question:

RQ1: How can we generate counterfactual explanations specifically for tree-based models?

Existing methods for generating counterfactual explanations for tree-based models are either based on heuristics or on integer linear programming techniques. The former do not necessarily converge to an optimal solution, while the latter can be extremely computationally intensive.

We answer **RQ1** by generating probabilistic approximations of tree-based models, which are differentiable and can therefore be used within a standard gradient-based optimization framework. Our experimental results show that our algorithm can generate minimal counterfactual explanations in an efficient and reliable manner.

2.1 INTRODUCTION

As Machine Learning (ML) models are prominently applied and their outcomes have a substantial effect on the general population, there is an increased demand for understanding what contributes to their predictions [Doshi-Velez and Kim, 2018]. For an individual who is affected by the predictions of these models, it would be useful to have an *actionable* explanation – one that provides insight into how these decisions can be *changed*. The General Data Protection Regulation (GDPR) is an example of recently enforced regulation in Europe which gives an individual the right to an explanation for algorithmic decisions, making the interpretability problem a crucial one for organizations that wish to adopt more data-driven decision-making processes [EU, 2016].

Counterfactual explanations are a natural solution to this problem since they frame the explanation in terms of what input (feature) changes are required to change the

This chapter was published at the AAAI Conference on Artificial Intelligence (AAAI 2022) under the title "FOCUS: Flexible Optimizable Counterfactual Explanations for Tree Ensembles" [Lucic et al., 2022b].

output (prediction). For instance, a user may be denied a loan based on the prediction of an ML model used by their bank. A counterfactual explanation could be: "Had your income been $\in 1000$ higher, you would have been approved for the loan." We focus on finding optimal counterfactual explanations: the minimal changes to the input required to change the outcome.

Counterfactual explanations are based on *counterfactual examples*: generated instances that are close to an existing instance but have an alternative prediction. The difference between the original instance and the counterfactual example is the counterfactual explanation. Wachter et al. [2018] propose framing the problem as an optimization task, but their work assumes that the underlying machine learning models are differentiable, which excludes an important class of widely applied and highly effective non-differentiable models: tree ensembles. We propose a method that relaxes this assumption and builds upon the work of Wachter et al. by introducing differentiable approximations of tree ensembles that can be used in such an optimization framework. Alternative non-optimization approaches for generating counterfactual explanations for tree ensembles involve an extensive search over many possible paths in the ensemble that could lead to an alternative prediction [Tolomei et al., 2017].

Given a trained tree-based model f, we probabilistically approximate f by replacing each split in each tree with a sigmoid function centred at the splitting threshold. If f is an ensemble of trees, then we also replace the maximum operator with a softmax. This approximation allows us to generate a counterfactual example \bar{x} for an instance x based on the minimal perturbation of x such that the prediction changes: $y_x \neq y_{\bar{x}}$, where y_x and $y_{\bar{x}}$ are the labels f assigns to x and \bar{x} , respectively. This leads us to our main research question in this chapter:

Are counterfactual examples generated by our method closer to the original input instances than those generated by existing heuristic methods?

Our main findings are that our proposed method is (i) a more *effective* counterfactual explanation method for tree ensembles than previous approaches since it manages to produce counterfactual examples that are closer to the original input instances than existing approaches; (ii) a more *efficient* counterfactual explanation method for tree ensembles since it is able to handle larger models than existing approaches; and (iii) a more *reliable* counterfactual explanation method for tree ensembles since it is able to generate counterfactual explanations for all instances in a dataset, unlike existing approaches specific to tree ensembles.

In the following sections, we examine existing work related to ours (Section 2.2) and formalize the counterfactual explanation problem (Section 2.3). We then describe the details of our method, Flexible Optimizable CoUnterfactual Explanations for Tree EnsembleS (FOCUS), in Section 2.4. In Section 2.5, we explain the experimental setup, followed by the experimental results in Sections 2.6 and 2.7. We analyze our findings in Section 2.8 and conclude in Section 2.9.

2.2 RELATED WORK

Our work is related to counterfactual explanations in general (Section 2.2.1), algorithmic recourse (Section 2.2.2), adversarial examples (Section 2.2.3), and differentiable tree-based models (Section 2.2.4).

2.2.1 Counterfactual Explanations

Counterfactual examples have been used in a variety of ML areas, such as reinforcement learning [Madumal et al., 2019], deep learning [Alaa et al., 2017], and explainable AI (XAI). Previous XAI methods for generating counterfactual examples are either modelagnostic [Poyiadzi et al., 2020; Karimi et al., 2020a; Laugel et al., 2018; Van Looveren and Klaise, 2021; Mothilal et al., 2020] or model-specific [Wachter et al., 2018; Grath et al., 2018; Tolomei et al., 2017; Kanamori et al., 2020; Russell, 2019; Dhurandhar et al., 2018]. Model-agnostic approaches treat the original model as a "black-box" and only assume query access to the model, whereas model-specific approaches typically do not make this assumption and can therefore make use of its inner workings.

Our work is a model-specific approach for generating counterfactual examples through optimization. Previous model-specific work for generating counterfactual examples through optimization has solely been conducted on differentiable models [Wachter et al., 2018; Grath et al., 2018; Dhurandhar et al., 2018].

2.2.2 Algorithmic Recourse

Algorithmic recourse is a line of research that is closely related to counterfactual explanations, except that methods for algorithmic recourse include the additional restriction that the resulting explanation must be *actionable* [Ustun et al., 2019; Joshi et al., 2020; Karimi et al., 2021, 2020c]. This is done by selecting a subset of the features to which perturbations can be applied in order to avoid explanations that suggest impossible or unrealistic changes to the feature values (i.e., change *age* from $50 \rightarrow 25$ or change *marital_status* from MARRIED \rightarrow UNMARRIED). Although this work has produced impressive theoretical results, it is unclear how realistic they are in practice, especially for complex ML models such as tree ensembles. Existing algorithmic recourse methods cannot solve our task because they (i) are either restricted to solely linear [Ustun et al., 2019] or differentiable [Joshi et al., 2020] models, or (ii) require access to causal information [Karimi et al., 2021, 2020c], which is rarely available in real world settings.

2.2.3 Adversarial Examples

Adversarial examples are a type of counterfactual example with the additional constraint that the minimal perturbation results in an alternative prediction that is *incorrect*. There are a variety of methods for generating adversarial examples [Goodfellow et al., 2015; Szegedy et al., 2014; Su et al., 2019; Brown et al., 2018]; a more complete overview can be found in the work of Biggio and Roli [2018]. The main difference between adversarial examples and counterfactual examples is in the intent: adversarial examples

are meant to *fool* the model, whereas counterfactual examples are meant to *explain* the model.

2.2.4 Differentiable Tree-based Models

Part of our contribution involves constructing differentiable versions of tree ensembles by replacing each splitting threshold with a sigmoid function. This can be seen as using a (small) neural network to obtain a smooth approximation of each tree. Neural decision trees [Balestriero, 2017; Yang et al., 2018] are also differentiable versions of trees, which use a full neural network instead of a simple sigmoid. However, these do not optimize for approximating an already trained model. Therefore, unlike our method, they are not an obvious choice for finding counterfactual examples for an existing model. Soft decision trees [Hinton et al., 2014] are another example of differentiable trees, which instead approximate a neural network with a decision tree. This can be seen as the inverse of our task.

2.3 PROBLEM FORMULATION

A counterfactual explanation for an instance x and a model f, Δ_x , is a minimal perturbation of x that changes the prediction of f. f is a probabilistic classifier, where $f(y \mid x)$ is the probability of x belonging to class y according to f. The prediction of f for x is the most probable class label $y_x = \arg\max_y f(y \mid x)$, and a perturbation \bar{x} is a counterfactual example for x if, and only if, $y_x \neq y_{\bar{x}}$, that is:

$$\arg\max_{y} f(y \mid x) \neq \arg\max_{y'} f(y' \mid \bar{x}). \tag{2.1}$$

In addition to changing the prediction, the distance between x and \bar{x} should also be minimized. We therefore define an *optimal counterfactual example* \bar{x}^* as:

$$\bar{x}^* := \arg\min_{\bar{x}} d(x, \bar{x}) \text{ such that } y_x \neq y_{\bar{x}},$$
 (2.2)

where $d(x, \bar{x})$ is a differentiable distance function. The corresponding *optimal counter-factual explanation* Δ_x^* is:

$$\Delta_x^* = \bar{x}^* - x. \tag{2.3}$$

This definition aligns with previous ML work on counterfactual explanations [Laugel et al., 2018; Karimi et al., 2020a; Tolomei et al., 2017]. We note that this notion of *optimality* is purely from an algorithmic perspective and does not necessarily translate to optimal changes in the real world, since the latter are completely dependent on the context in which they are applied. It should be noted that if the loss space is non-convex, it is possible that more than one optimal counterfactual explanation exists.

Minimizing the distance between x and \bar{x} should ensure that \bar{x} is as close to the decision boundary as possible. This distance indicates the effort it takes to apply the perturbation in practice, and an optimal counterfactual explanation shows how a

prediction can be changed with the least amount of effort. An optimal explanation provides the user with interpretable and potentially actionable feedback related to understanding the predictions of model f.

Wachter et al. [2018] recognized that counterfactual examples can be found through gradient descent if the task is cast as an optimization problem. Specifically, they use a loss consisting of two components: (i) a prediction loss to change the prediction of $f: \mathcal{L}_{pred}(x, \bar{x} \mid f)$, and (ii) a distance loss to minimize the distance $d: \mathcal{L}_{dist}(x, \bar{x} \mid d)$. The complete loss is a linear combination of these two parts, with a weight $\beta \in \mathbb{R}_{>0}$:

$$\mathcal{L}(x, \bar{x} \mid f, d) = \mathcal{L}_{pred}(x, \bar{x} \mid f) + \beta \mathcal{L}_{dist}(x, \bar{x} \mid d). \tag{2.4}$$

The assumption here is that an optimal counterfactual example \bar{x}^* can be found by minimizing the overall loss:

$$\bar{x}^* = \arg\min_{\bar{x}} \mathcal{L}(x, \bar{x} \mid f, d). \tag{2.5}$$

Wachter et al. [2018] propose a prediction loss \mathcal{L}_{pred} based on the mean-squared-error. A clear limitation of this approach is that it assumes f is differentiable. This excludes many commonly used ML models, including tree-based models, on which we focus in this work.

2.4 METHOD: FOCUS

To mimic many real-world scenarios, we assume there exists a trained model f that we need to explain. The goal here is not to create a new, inherently interpretable tree-based model, but rather to explain a model that already exists.

2.4.1 Loss Function Definitions

We use a hinge-loss since we assume a classification task:

$$\mathcal{L}_{pred}(x, \bar{x} \mid f) = \mathbb{1} \left[\arg \max_{y} f(y \mid x) = \arg \max_{y'} f(y' \mid \bar{x}) \right] \cdot f(y' \mid \bar{x}). \tag{2.6}$$

Allowing for flexibility in the choice of distance function allows us to tailor the explanations to the end-users' needs. We make the preferred notion of *minimality* explicit through the choice of distance function. Given a differentiable distance function *d*, the distance loss is:

$$\mathcal{L}_{dist}(x,\bar{x}) = d(x,\bar{x}). \tag{2.7}$$

Building off of Wachter et al. [2018], we propose incorporating differentiable approximations of non-differentiable models to use in the gradient-based optimization framework. Since the approximation \tilde{f} is derived from the original model f, it should match f closely: $\tilde{f}(y\mid x)\approx f(y\mid x)$. We define the approximate prediction loss as follows:

$$\widetilde{\mathcal{L}}_{pred}(x, \bar{x} \mid f, \tilde{f}) = \mathbb{1} \left[\arg \max_{y} f(y \mid x) = \arg \max_{y'} f(y' \mid \bar{x}) \right] \cdot \widetilde{f}(y' \mid \bar{x}). \quad (2.8)$$

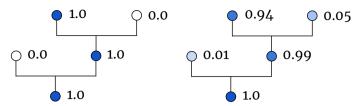


Figure 2.1: Left: A decision tree \mathcal{T} and node activations for a single instance. Right: a differentiable approximation of the same tree $\widetilde{\mathcal{T}}$ and activations for the same instance.

This loss is based both on the original model f and the approximation \tilde{f} : the loss is active as long as the prediction according to f has not changed, but its gradient is based on the differentiable \tilde{f} . This prediction loss encourages the perturbation to have a different prediction than the original instance by penalizing an unchanged instance. The approximation of the complete loss becomes:

$$\widetilde{\mathcal{L}}(x,\bar{x}\mid f,\tilde{f},d) = \widetilde{\mathcal{L}}_{pred}(x,\bar{x}\mid f,\tilde{f}) + \beta \cdot \mathcal{L}_{dist}(x,\bar{x}\mid d). \tag{2.9}$$

Since we assume that it approximates the complete loss,

$$\widetilde{\mathcal{L}}(x,\bar{x}\mid f,\tilde{f},d) \approx \mathcal{L}(x,\bar{x}\mid f,d),$$
 (2.10)

we also assume that an optimal counterfactual example can be found by minimizing it:

$$\bar{x}^* \approx \arg\min_{\bar{x}} \ \widetilde{\mathcal{L}}(x, \bar{x} \mid f, \tilde{f}, d).$$
 (2.11)

2.4.2 Tree-based Models

To obtain the differentiable approximation \tilde{f} of f, we construct a probabilistic approximation of the original tree ensemble f. Tree ensembles are based on decision trees; a single decision tree \mathcal{T} uses a binary-tree structure to make predictions about an instance x based on its features. Figure 2.1 shows a simple decision tree consisting of five nodes. A node j is activated if its parent node p_j is activated and feature x_{f_j} is on the correct side of the threshold θ_j ; which side is the correct side depends on whether j is a left or right child. The root note is an exception, it is always activated. Let $t_j(x)$ indicate if node j is activated:

$$t_j(x) = \begin{cases} 1, & \text{if } j \text{ is the root,} \\ t_{p_j}(x) \cdot \mathbb{1}[x_{f_j} > \theta_j], & \text{if } j \text{ is a left child,} \\ t_{p_j}(x) \cdot \mathbb{1}[x_{f_j} \leq \theta_j], & \text{if } j \text{ is a right child.} \end{cases}$$
 (2.12)

 $\forall x, t_0(x) = 1$. Nodes that have no children are called *leaf nodes*; an instance x always ends up in a single leaf node. Every leaf node j has its own predicted distribution $\mathcal{T}(y \mid j)$; the prediction of the full tree is given by its activated leaf node. Let \mathcal{T}_{leaf} be the set of leaf nodes in \mathcal{T} , then:

$$(j \in \mathcal{T}_{leaf} \wedge t_j(x) = 1) \to \mathcal{T}(y \mid x) = \mathcal{T}(y \mid j).$$
 (2.13)

Alternatively, we can reformulate this as a sum over leaves:

$$\mathcal{T}(y \mid x) = \sum_{j \in \mathcal{T}_{leaf}} t_j(x) \cdot \mathcal{T}(y \mid j). \tag{2.14}$$

Generally, tree ensembles are deterministic; let f be an ensemble of M many trees with weights $\omega_m \in \mathbb{R}$, then:

$$f(y \mid x) = \arg\max_{y'} \sum_{m=1}^{M} \omega_m \cdot \mathcal{T}_m(y' \mid x).$$
 (2.15)

2.4.3 Approximations of Tree-based Models

If f is not differentiable, we are unable to calculate its gradient with respect to the input x. However, the non-differentiable operations in our formulation are (i) the indicator function, and (ii) a maximum operation, both of which can be approximated by differentiable functions. First, we introduce the $\widetilde{t}_j(x)$ function that approximates the activation of node $j\colon \widetilde{t}_j(x)\approx t_j(x)$, using a sigmoid function with parameter $\sigma\in\mathbb{R}_{>0}$: $sig(z)=(1+\exp(\sigma\cdot z))^{-1}$ and

$$\widetilde{t}_{j}(x) = \begin{cases} 1, & \text{if } j \text{ is the root,} \\ \widetilde{t}_{p_{j}}(x) \cdot \operatorname{sig}(\theta_{j} - x_{f_{j}}), & \text{if } j \text{ is left child,} \\ \widetilde{t}_{p_{j}}(x) \cdot \operatorname{sig}(x_{f_{j}} - \theta_{j}), & \text{if } j \text{ is right child.} \end{cases}$$
 (2.16)

As σ increases, $\widetilde{t_j}$ approximates t_j more closely. Next, we introduce a *tree approximation*:

$$\widetilde{\mathcal{T}}(y \mid x) = \sum_{j \in \mathcal{T}_{leaf}} \widetilde{t}_j(x) \cdot \mathcal{T}(y \mid j). \tag{2.17}$$

The approximation $\widetilde{\mathcal{T}}$ uses the same tree structure and thresholds as \mathcal{T} . However, its activations are no longer deterministic but instead are dependent on the distance between the feature values x_{f_j} and the thresholds θ_j . Lastly, we replace the maximum operation of f by a softmax with temperature $\tau \in \mathbb{R}_{>0}$, resulting in:

$$\tilde{f}(y \mid x) = \frac{\exp\left(\tau \cdot \sum_{m=1}^{M} \omega_m \cdot \widetilde{\mathcal{T}}_m(y \mid x)\right)}{\sum_{y'} \exp\left(\tau \cdot \sum_{m=1}^{M} \omega_m \cdot \widetilde{\mathcal{T}}_m(y' \mid x)\right)}.$$
(2.18)

The approximation \tilde{f} is based on the original model f and the parameters σ and τ . This approximation is applicable to any tree-based model, and how well \tilde{f} approximates f depends on the choice of σ and τ . The approximation is potentially perfect since

$$\lim_{\substack{x \to \infty \\ x \to \infty}} \tilde{f}(y \mid x) = f(y \mid x). \tag{2.19}$$

2.4.4 Our Method: FOCUS

We call our method FOCUS: Flexible Optimizable CounterfactUal Explanations for Tree EnsembleS. It takes as input an instance x, a tree-based classifier f, and two hyperparameters: σ and τ , which we use to create the approximation \tilde{f} . Following Equation 2.11, FOCUS outputs the optimal counterfactual example \bar{x}^* , from which we derive the optimal counterfactual explanation $\Delta_x^* = \bar{x}^* - x$.

2.4.5 Effects of Hyperparameters

Increasing σ in \tilde{f} eventually leads to exact approximations of the indicator functions, while increasing τ in \tilde{f} leads to a completely unimodal softmax distribution. It should be noted that our approximation \tilde{f} is not intended to replace the original model f but rather to create a differentiable version of f from which we can generate counterfactual examples through optimization. In practice, the original model f would still be used to make predictions and the approximation would solely be used to generate counterfactual examples.

2.5 EXPERIMENTAL SETUP

We consider 42 experimental settings to find the best counterfactual explanations using FOCUS. We jointly tune the hyperparameters of FOCUS (σ , τ , β , α) using Adam [Kingma and Ba, 2015] for 1,000 iterations. We choose the hyperparameters that produce (i) a valid counterfactual example for every instance in the dataset, and (ii) the smallest mean distance between corresponding pairs (x, \bar{x}).

We evaluate FOCUS on four binary classification datasets: *Wine Quality* [UCI, 2009], *HELOC* [FICO, 2017], *COMPAS* [Ofer, 2017], and *Shopping* [UCI, 2019]. For each dataset, we train three types of tree-based models: Decision Trees (DT), Random Forests (RF), and Adaptive Boosting Trees (AB) with DTs as the base learners. We compare against two baselines that generate counterfactual examples for tree ensembles based on the inner workings of the model: Feature Tweaking (FT) by Tolomei et al. [2017] and Distribution-Aware Counterfactual Explanations (DACE) by Kanamori et al. [2020].

2.5.1 Datasets

We evaluate FOCUS on four binary classification tasks using the following datasets:

- The Wine Quality dataset [UCI, 2009] has 4,898 instances and 11 features. The task is about predicting the quality of white wine on a 0–10 scale. We adapt this to a binary classification setting by labelling the wine as "high quality" if the quality is ≥ 7 .
- The *HELOC* dataset [FICO, 2017] has 10,459 instances and 23 features. The task is from the Explainable Machine Learning Challenge at NeurIPS 2017, where the task is to predict whether or not a customer will default on their loan.

- The *COMPAS* dataset [Ofer, 2017] has 6,172 instances and 6 features. It is used for detecting bias in ML systems, where the task is predicting whether or not a criminal defendant will reoffend upon release.
- The *Shopping* dataset [UCI, 2019] has 12,330 instances and 9 features. The task entails predicting whether or not an online website visit results in a purchase.

We scale all features such that their values are in the range $\left[0,1\right]$ and remove categorical features

2.5.2 Models

We train three types of tree-based models on 70% of each dataset: Decision Trees (DTs), Random Forests (RFs), and Adaptive Boosting (AB) with DTs as the base learners. We use the remaining 30% to find counterfactual examples for this test set. In total we have 12 models (4 datasets \times 3 tree-based models).

2.5.3 Distance Functions

In our experiments, we generate different types of counterfactual explanations using different types of distance functions. We note that the flexibility of FOCUS allows for the use of any differentiable distance function. Euclidean distance measures the geometric displacement:

$$d_{Euclidean}(x,\bar{x}) = \sqrt{\sum_{i} (x_i - \bar{x}_i)^2}.$$
 (2.20)

Cosine distance measures the angle by which \bar{x} deviates from x – whether \bar{x} preserves the relationship between features in x:

$$d_{Cosine}(x,\bar{x}) = 1 - \frac{\sum_{i} (x_i \cdot \bar{x}_i)}{\|x\| \|\bar{x}\|}.$$
 (2.21)

Manhattan distance (i.e., L1-norm) measures per feature differences, minimizing the number of features perturbed and therefore inducing sparsity:

$$d_{Manhattan}(x,\bar{x}) = \sum_{i} |x_i - \bar{x}_i|. \tag{2.22}$$

When comparing against DACE [Kanamori et al., 2020], we use the Mahalanobis distance, since this is the distance function used in their novel cost function (see Equation 2.27):

$$d_{Mahalanobis}(x,\bar{x}|C) = \sqrt{(x-\bar{x})C^{-1}(x-\bar{x})}.$$
 (2.23)

C is the covariance matrix of x and \bar{x} , which allows us to account for correlations between features. When all features are uncorrelated, the Mahalanobis distance is equal to the Euclidean distance.

2.5.4 Evaluation Metrics

We evaluate the counterfactual examples produced by FOCUS based on how close they are to the original input using three metrics, in terms of four distance functions (see Section 2.5.3). The first evaluation metric is distance from the original input averaged over all examples, d_{mean} . Let X be the set of N original instances and \bar{X} be the corresponding set of N generated counterfactual examples. The *mean distance* is defined as:

$$d_{mean}(X, \bar{X}) = \frac{1}{N} \sum_{n=1}^{N} d(x^{(n)}, \bar{x}^{(n)}).$$
 (2.24)

The second evaluation metric is mean relative distance from the original input, d_{Rmean} . This metric helps us interpret individual improvements over the baselines; if $d_{Rmean} < 1$, FOCUS's counterfactual examples are on average closer to the original input compared to the baseline. Let \bar{X} be the set of counterfactual examples produced by FOCUS and let \bar{X}' be the set of counterfactual examples produced by a baseline. Then the mean relative distance is defined as:

$$d_{Rmean}(\bar{X}, \bar{X}') = \frac{1}{N} \sum_{n=1}^{N} \frac{d(x^{(n)}, \bar{x}^{(n)})}{d(x^{(n)}, \bar{x}'^{(n)})}.$$
 (2.25)

The third evaluation metric is the proportion of FOCUS's counterfactual examples that are closer to the original input in comparison to the baselines. For d we consider the Euclidean, Cosine, Manhattan, and Mahalanobis distance functions.

2.6 EXPERIMENT 1: FOCUS VS. FT

We compare FOCUS to the Feature Tweaking (FT) method by Tolomei et al. [2017] in terms of the evaluation metrics in Section 2.5.4. We consider 36 experimental settings (4 datasets \times 3 tree-based models \times 3 distance functions) when comparing FOCUS to FT. The results are listed in Table 2.1.

2.6.1 Baseline: Feature Tweaking

Feature Tweaking identifies the leaf nodes where the prediction of the leaf nodes do not match the original prediction y_x : it recognizes the set of leaves that if activated, $t_i(\bar{x}) = 1$, would change the prediction of a tree \mathcal{T} :

$$\mathcal{T}_{change} = \left\{ j \mid j \in \mathcal{T}_{leaf} \land y_x \neq \arg\max_{y} T(y \mid j) \right\}. \tag{2.26}$$

For every \mathcal{T} in f, Feature Tweaking (FT) generates a perturbed example per node in \mathcal{T}_{change} so that it is activated with at least an ϵ difference per threshold, and then selects the most optimal example (i.e., the one closest to the original instance). For every feature threshold θ_j involved, the corresponding feature is perturbed accordingly: $\bar{x}_{f_j} = \theta_j \pm \epsilon$. The result is a perturbed example that was changed minimally to activate a leaf node in \mathcal{T}_{change} . In our experiments, we test $\epsilon \in \{0.001, 0.005, 0.01, 0.1\}$, and

choose the ϵ that minimizes the mean distance to the original input, while maximizing the number of counterfactual examples generated.

The main problem with FT is that the perturbed examples are not necessarily counterfactual examples, since changing the prediction of a single tree \mathcal{T} does not guarantee a change in the prediction of the full ensemble f. Figure 2.2 shows all three perturbed examples generated by FT for a single instance. In this case, none of the generated examples change the model prediction and therefore none are valid counterfactual examples.

Figure 2.2 shows how FOCUS and FT handle an adaptive boosting ensemble using a two-feature ensemble with three trees. On the left is the decision boundary for a standard tree ensemble; the middle visualizes the positive leaf nodes that form the decision boundary; on the right is the approximated loss $\widetilde{\mathcal{L}}_{pred}$ and its gradient w.r.t. \bar{x} . The gradients push features close to thresholds harder and in the direction of the decision boundary if $\widetilde{\mathcal{L}}$ is convex.

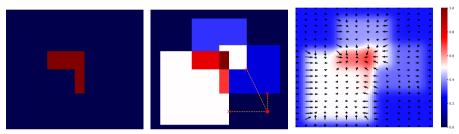


Figure 2.2: An example of how the FT baseline method (explained in Section 2.6.1) and our method handle an adaptive boosting ensemble with three trees. Left: decision boundary of the ensemble. Middle: three positive leaves that form the decision boundary, an example instance, and the perturbed examples suggested by FT. Right: approximated loss $\widetilde{\mathcal{L}}_{pred}$ and its gradient w.r.t. \bar{x} . The FT perturbed examples do not change the prediction of the forest, whereas the gradient of the differentiable approximation leads toward the true decision boundary.

2.6.2 Results

In terms of d_{mean} , FOCUS outperforms FT in 20 settings while FT outperforms FOCUS in 8 settings. The difference in d_{mean} is not significant in the remaining 8 settings. In general, FOCUS outperforms FT in settings using Euclidean and Cosine distance because in each iteration, FOCUS perturbs many of the features by a small amount. Since FT perturbs only the features associated with an individual leaf, we expected that it would perform better for Manhattan distance but our results show that this is not the case.

Table 2.1: Evaluation results for Experiment 1 comparing FOCUS and FT counterfactual examples. Significant improvements and losses over the baseline (FT) are denoted by \P and \P , respectively (p < 0.05, two-tailed t-test,); $^{\circ}$ denotes no significant difference; $^{\otimes}$ denotes settings where the baseline cannot find a counterfactual example for every instance.

			E	uclideaı	1		Cosine		M	anhatta	n
Dataset	Metric	Method	DT	RF	AB	DT	RF	AB	DT	RF	AB
Wine	d_{mean}	FT FOCUS	0.269 0.268 °	0.174 0.188▲	0.267 [⊗] 0.188 [▼]	0.030 0.003 ▼	0.017 0.008 ▼	0.034 [⊗] 0.014 [▼]	0.269 0.268 °	0.223 0.312 [▲]	0.382 [⊗] 0.360 [▼]
Quality	$d_{Rmean} \ \%_{closer}$	FOCUS/FT FOCUS < FT	0.990 100%	1.256 21.0%	0.649 87.5%	0.066 100%	0.821 80.8%	0.312 95.1%	0.990 100%	1.977 5.4%	0.924 58.6%
HELOC	d_{mean}	FT FOCUS	0.120 0.133▲	0.210 0.186 ♥	0.185 0.136 ♥	0.003 0.001 ▼	0.008 0.002 ▼	0.007 0.001 ▼	0.135 0.152▲	0.278 0.284°	0.198 0.203°
	$d_{Rmean} \ \%_{closer}$	FOCUS/FT FOCUS < FT	1.169 16.6%	0.942 57.9%	0.907 71.9%	0.303 91.6%	0.285 91.5%	0.421 92.9%	1.252 51.3%	1.144 43.6%	1.364 24.2%
COMPAS	d_{mean}	FT FOCUS	0.082 0.092▲	0.075 0.079°	0.081 0.076 ♥	0.013 0.008 ▼	0.014 0.011 ▼	0.015 0.007 ▼	0.086 0.093▲	0.078 0.085°	0.085 0.090°
	$d_{Rmean} \ \%_{closer}$	FOCUS/FT FOCUS < FT	1.162 29.4%	1.150 22.6%	1.062 44.8%	0.473 82.7%	0.965 68.0%	0.539 84.8%	1.182 65.8%	1.236 36.2%	1.155 66.9%
Shopping	d_{mean}	FT FOCUS	0.119 0.142 [▲]	0.028 0.025 ▼	0.126 [⊗] 0.028 [▼]	0.050 0.055 [▲]	0.027 0.013 ♥	0.131 [⊗] 0.006 [▼]	0.121 0.128°	0.030 0.026 ♥	0.142 [⊗] 0.046 [▼]
	$d_{Rmean} \ \%_{closer}$	FOCUS/FT FOCUS < FT	1.051 40.2%	1.053 36.1%	0.218 99.6%	0.795 44.4%	0.482 86.1%	0.074 99.5%	0.944 55.8%	0.796 81.9%	0.312 97.1%

There is no clear winner between FT and FOCUS for Manhattan distance. We also see that FOCUS usually outperforms FT in settings using Random Forests (RF) and Adaptive Boosting (AB), while the opposite is true for Decision Trees (DT).

Overall, we find that FOCUS is effective and efficient for finding counterfactual explanations for tree-based models. Unlike the FT baseline, FOCUS finds valid counterfactual explanations for *every* instance across all settings. In the majority of tested settings, FOCUS's explanations are substantial improvements in terms of distance to the original inputs, across all three metrics.

2.7 EXPERIMENT 2: FOCUS VS. DACE

The flexibility of FOCUS allows us to plug in our choice of differentiable distance function. To compare against DACE [Kanamori et al., 2020], we use the Mahalanobis distance for both (i) generation of FOCUS explanations, and (ii) evaluation in comparison to DACE, since this is the distance function used in the DACE loss function (see Equation 2.27 in Section 2.7.1).

We found two main limitations of DACE: (i) in all of our settings, it can only generate counterfactual examples for a subset of the test set, and (ii) it is limited by the size of the tree-based model. All hyperparameter settings are listed in the Appendix to this chapter.

2.7.1 Baseline: DACE

DACE generates counterfactual examples that account for the underlying data distribution through a novel cost function using Mahalanobis distance and a local outlier factor (LOF):

$$d_{DACE}(x,\bar{x}|X,C) = d_{Mahalanobis}^{2}(x,\bar{x}|C) + \lambda q_{k}(x,\bar{x}|X), \qquad (2.27)$$

where C is the covariance matrix, q_k is the k-LOF [Breunig et al., 2000], X is the training set, and λ is the trade-off parameter. The k-LOF measures the degree to which an instance is an outlier in the context of its k-nearest neighbors. To generate counterfactual examples, DACE formulates the task as a mixed-integer linear optimization problem and uses the CPLEX Optimizer to solve it. We refer the reader to the original paper for a more detailed overview of this cost function. The q_k term in the loss function penalizes counterfactual examples that are outliers, and therefore decreasing λ results in a greater number of counterfactual examples. In our experiments, we test $\lambda \in \{0.001, 0.01, 0.1, 0.5, 1.0\}$, and choose the λ that minimizes the mean distance to the original input, while maximizing the number of counterfactual examples generated.

We were only able to run DACE on 6 out of our 12 models because the problem size is too large (i.e., there are too many model parameters for DACE) for the remaining

¹We use k = 1 in our experiments, since this is the value of k that is supported in the code kindly provided to us by the authors, for which we are very grateful.

²http://www.ibm.com/analytics/cplex-optimizer

6 models when using the free Python API of CPLEX (the optimizer used in DACE). Specifically, we were unable to run DACE on the following settings:

- Wine Quality AB (100 trees, max depth 4)
- Wine Quality RF (500 trees, max depth 4)
- HELOC RF (500 trees, max depth 4)
- HELOC AB (100 trees, max depth 8)
- COMPAS RF (500 trees, max depth 4)
- Shopping RF (500 trees, max depth 8).

Therefore, when comparing against DACE, we have 6 experimental settings (6 models \times 1 distance function). We note that these are not unreasonable model sizes, and that unlike DACE, FOCUS can be applied to all 12 models (see Table 2.1).

2.7.2 Results

Table 2 shows the results for the 6 settings we could run DACE on. We were only able to run DACE on 6 out of our 12 models because the problem size is too large (i.e., DACE has too many model parameters) for the remaining 6 models when using the free Python API of CPLEX (the optimizer used in DACE). Therefore, when comparing against DACE, we have 6 experimental settings (6 models \times 1 distance function).

We found that DACE can only generate counterfactual examples for a small subset of the test set, regardless of the λ -value, as opposed to FOCUS, which can generate counterfactual examples for the entire test set in all cases. To compute d_{mean} , d_{Rmean} , and $\%_{closer}$, we compare FOCUS and DACE only on the instances for which DACE was able to generate a counterfactual example. We find that FOCUS significantly outperforms DACE in 5 out of 6 settings in terms of all three evaluation metrics, indicating that FOCUS explanations are indeed more minimal than those produced by DACE. FOCUS is also more reliable since (i) it is not restricted by model size, and (ii) it can generate counterfactual examples for all instances in the test set.

2.8 DISCUSSION AND ANALYSIS

Figure 2.3 shows the mean Manhattan distance of the perturbed examples in each iteration of FOCUS, along with the proportion of perturbations resulting in valid counterfactual examples found for two datasets (we omit the others due to space considerations). These trends are indicative of all settings: the mean distance increases until a counterfactual example has been found for every x, after which the mean distance starts to decrease. This seems to be a result of the hinge-loss in FOCUS, which first prioritizes finding a valid counterfactual example (see Equation 2.1), then decreasing the distance between x and \bar{x} .

Table 2.2: Evaluation results for Experiment 2 comparing FOCUS and DACE counterfactual examples in terms of Mahalanobis distance. Significant improvements over the baseline are denoted by \P (p < 0.05, two-tailed t-test,). $^{\circ}$ denotes no significant difference.

		Wine	HELOC	COMPAS		Shopping	
Metric	Method	DT	DT	DT	AB	DT	AB
d_{mean}	DACE FOCUS	1.325 0.542 ▼	1.427 0.810	0.814 ′0.776 °	1.570 0.636	0.050 0.023	3.230 0.303
d_{Rmean}	FOCUS / DACE	0.420	0.622	1.18	0.372	0.449	0.380
$\%_{closer}$	FOCUS < DACE	100%	94.5%	29.9%	96.1%	99.4%	90.8%
# CFs	DACE	241	1342	842	700	362	448
found	FOCUS	1470	3138	1852	1852	3699	3699
# obs in	dataset	1470	3138	1852	1852	3699	3699

2.8.1 Case Study: Credit Risk

As a practical example, we investigate what FOCUS explanations look like for individuals in the HELOC dataset. Here, the task is to predict whether or not an individual will default on their loan. This has consequences for loan approval: individuals who are predicted as defaulting will be denied a loan. For these individuals, we want to understand how they can change their profile such that they are approved. Given an individual who has been denied a loan from a bank, a counterfactual explanation could be:

Your loan application has been denied. In order to have your loan application approved, you need to (i) increase your ExternalRiskEstimate score by 62, and (ii) decrease your NetFractionRevolvingBurden by 58.

Figure 2.4 shows four counterfactual explanations generated using different distance functions for the same individual and same model. We see that the Manhattan explanation only requires a few changes to the individual's profile, but the changes are large. In contrast, the individual changes in the Euclidean explanation are smaller but there are more of them. In settings where there are significant dependencies between features, the Cosine explanations may be preferred since they are based on perturbations that try to preserve the relationship between features. For instance, in the *Wine Quality* dataset, it would be difficult to change the amount of citric acid without affecting the pH level. The Mahalanobis explanations would be useful when it is important to take into account not only correlations between features, but also the training data distribution. This flexibility allows users to choose what kind of explanation is best suited for their problem.

Different distance functions can result in different *magnitudes* of feature perturbations as well as different *directions*. For example, the Cosine explanation suggests

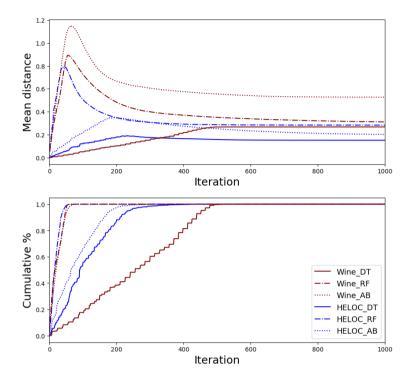


Figure 2.3: Mean distance (top) and cumulative % (bottom) of counterfactual examples in each iteration of FOCUS for Manhattan explanations.

increasing *PercentTradesWBalance*, while the Mahalanobis explanations suggests decreasing it. This is because the loss space of the underlying RF model is highly non-convex, and therefore there is more than one way to obtain an alternative prediction. When using complex models such as tree ensembles, there are no monotonicity guarantees. In this case, both options result in valid counterfactual examples.

We examine the Manhattan explanation in more detail. We see that FOCUS suggests two main changes: (i) increasing the *ExternalRiskEstimate*, and (ii) decreasing the *NetFractionRevolvingBurden*. We obtain the definitions and expected trends from the data dictionary created by the authors of the dataset. The *ExternalRiskEstimate* is a "consolidated version of risk markers" (i.e., a credit score). A higher score is better: as one's *ExternalRiskEstimate* increases, the probability of default decreases. The *NetFractionRevolvingBurden* is the "revolving balance divided by the credit limit" (i.e., utilization). A lower value is better: as one's *NetFractionRevolvingBurden* increases, the probability of default increases. We find that the changes suggested by FOCUS are fairly consistent with the expected trends in the data dictionary, as opposed to suggesting nonsensical changes such as increasing one's utilization to decrease the probability of default.

Decreasing one's utilization is heavily dependent on the specific situation: an individual who only supports themselves might have more control over their spending

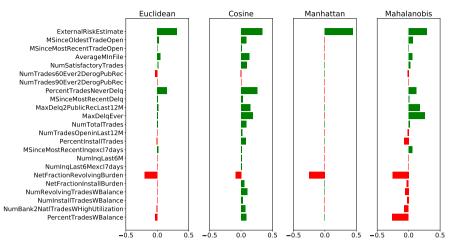


Figure 2.4: FOCUS explanations for the same model and same x based on different distance functions. Green and red indicate increases and decreases in feature values, respectively. Perturbation values are based on normalized feature values. Left: Euclidean explanation perturbs several features, but only slightly. Middle Left: Cosine explanation perturbs almost all of the features. Middle Right: Manhattan explanation perturbs two features substantially. Right: Mahalanobis explanation perturbs almost all of the features.

in comparison to someone who has multiple dependents. An individual can decrease their utilization in two ways: (i) decreasing their spending, or (ii) increasing their credit limit (or a combination of the two). We can postulate that (i) is more "actionable" than (ii), since (ii) is usually a decision made by a financial institution. However, the degree to which an individual can actually change their spending habits is completely dependent on their specific situation: an individual who only supports themselves might have more control over their spending than someone who has multiple dependents. In either case, we argue that deciding what is (not) actionable is not a decision for the developer to make, but for the individual who is affected by the decision. Counterfactual examples should be used as part of a human-in-the-loop system and not as a final solution.

The individual should know that utilization is an important component of the model, even if it is not necessarily "actionable" for them. We also note that it is unclear how exactly an individual would change their credit score without further insight into how the score was calculated (i.e., how the risk markers were consolidated). It should be noted that this is not a shortcoming of FOCUS, but rather of using features that are uninterpretable on their own, such as credit scores. Although FOCUS explanations cannot tell a user precisely how to increase their credit score, it is still important for the individual to know that their credit score is an important factor in determining their probability of getting a loan, as this empowers them to ask questions about how the score was calculated (i.e., how the risk markers were consolidated).

2.9 CONCLUSION

In this chapter, we propose an explanation method for tree-based classifiers, FOCUS, which casts the problem of finding counterfactual examples as a gradient-based optimization task and provides a differentiable approximation of tree-based models to be used in the optimization framework.

Given an input instance x, FOCUS generates an optimal counterfactual example based on the minimal perturbation to the input instance x which results in an alternative prediction from a model f. Unlike previous methods that assume the underlying classification model is differentiable, we propose a solution for when f is a non-differentiable, tree-based model that provides a differentiable approximation of f that can be used to find counterfactual examples using gradient-based optimization techniques. In the majority of experiments, examples generated by FOCUS are significantly closer to the original instances in terms of three different evaluation metrics compared to those generated by the baselines. FOCUS is able to generate valid counterfactual examples for all instances across all datasets, and the resulting explanations are flexible depending on the distance function.

This answers **RQ1**: we can generate counterfactual explanations for tree-based models using a combination of (i) gradient-based optimization techniques, and (ii) differentiable approximations of tree-based models which are used within the optimization framework. In the following chapter, we will investigate how to extend our method to accommodate different types of data such as graphs.

REPRODUCIBILITY

To facilitate the reproducibility of this work, our code is available at https://github.com/a-lucic/focus.

APPENDIX

Here we detail the FOCUS hyperparameters across the 42 settings in Experiments 1 and 2: σ indicates the steepness of the sigmoid function in Equation 2.16; τ is the temperature of the softmax in Equation 2.18; β is the trade-off parameter in Equation 2.9; α is the learning rate of Adam.

Table 2.3: FOCUS hyperparameters used in Experiment 1 comparing FOCUS with FT and RP using Euclidean distance.

Dataset	Model	Num Trees	Max Depth	σ	au	β	α
	DT	1	2	1	10	0.05	0.001
Wine	RF	500	4	10	2	0.05	0.005
	AB	100	4	5	1	0.05	0.005
	DT	1	4	2	10	0.05	0.001
HELOC	RF	500	4	10	5	0.05	0.005
	AB	100	8	10	1	0.05	0.001
	DT	1	4	6	10	0.05	0.005
COMPAS	RF	500	4	7	3	0.01	0.001
	AB	100	2	10	1	0.01	0.005
Shopping	DT	1	4	2	10	0.05	0.005
	RF	500	8	5	5	0.05	0.005
	AB	100	2	10	1	0.05	0.001

Table 2.4: FOCUS hyperparameters used in Experiment 1 comparing FOCUS with FT and RP using Cosine distance.

Dataset	Model	Num Trees	Max Depth	σ	au	β	α
	DT	1	2	1	10	0.05	0.005
Wine	RF	500	4	10	1	0.05	0.005
	AB	100	4	1	1	0.01	0.005
	DT	1	4	2	10	0.05	0.005
HELOC	RF	500	4	5	5	0.05	0.005
	AB	100	8	1	1	0.05	0.005
	DT	1	4	10	10	0.05	0.005
COMPAS	RF	500	4	10	6	0.01	0.005
	AB	100	2	10	1	0.05	0.005
Shopping	DT	1	4	10	10	0.05	0.001
	RF	500	8	1	1	0.01	0.001
	AB	100	2	10	5	0.05	0.001

Table 2.5: FOCUS hyperparameters used in Experiment 1 comparing FOCUS with FT and RP using Manhattan distance.

Dataset	Model	Num Trees	Max Depth	σ	au	β	α
	DT	1	2	1	10	0.05	0.001
Wine	RF	500	4	10	10	0.01	0.005
	AB	100	4	6	1	0.01	0.005
	DT	1	4	2	10	0.05	0.001
HELOC	RF	500	4	5	5	0.01	0.005
	AB	100	8	4	1	0.05	0.001
	DT	1	4	6	10	0.01	0.005
COMPAS	RF	500	4	4	1	0.05	0.001
	AB	100	2	5	10	0.05	0.005
Shopping	DT	1	4	2	10	0.05	0.005
	RF	500	8	10	1	0.05	0.001
	AB	100	2	10	1	0.05	0.001

Table 2.6: FOCUS hyperparameters used in Experiment 2 comparing FOCUS with DACE using Mahalanobis distance.

Dataset	Model	Num Trees	Max Depth	σ	au	β	α
Wine	DT	1	2	5	10	0.01	0.001
HELOC	DT	1	4	5	10	0.01	0.001
COMPAS	DT	1	4	5	10	0.01	0.005
	AB	100	2	4	2	0.005	0.001
Shopping	DT	1	4	4	10	0.01	0.005
	AB	100	2	10	1	0.01	0.001

Counterfactual Explanations for Graph Neural Networks

In the previous chapter, we developed a method for generating counterfactual explanations specific to tree-based models using gradient-based optimization techniques. In this chapter, we address the following research question:

RQ2: How can we extend our explanation method for tree-based models to graph-based models?

Most existing methods for explaining predictions from GNNs are based on retrieving a subgraph of the original graph that is most relevant for the prediction. This differs from the counterfactual explanation problem where the task is to find the minimal perturbation to the original graph such that the prediction changes. At the time of writing, there were no existing methods for generating *counterfactual* explanations for GNNs.

We answer **RQ2** by first extending the counterfactual explanation problem formalization to the graph data setting, then applying the same gradient-based optimization techniques as in the previous chapter. Our experimental results show that our algorithm can reliably generate minimal and accurate counterfactual explanations for GNNs.

3.1 INTRODUCTION

Advances in machine learning (ML) have led to breakthroughs in several areas of science and engineering, ranging from computer vision, to natural language processing, to conversational assistants. Parallel to the increased performance of ML systems, there is an increasing call for the "understandability" of ML models [Goebel et al., 2018]. Understanding *why* an ML model returns a certain output in response to a given input is important for a variety of reasons such as model debugging, aiding decison-making, or fulfilling legal requirements [EU, 2016]. Having certified methods for interpreting ML predictions will help enable their use across a variety of applications [Miller, 2019].

This chapter was published at the International Conference on Artificial Intelligence and Statistics (AISTATS 2022) under the title "CF-GNNExplainer: Counterfactual Explanations for Graph Neural Networks" [Lucic et al., 2022c].

Explainable AI (XAI) refers to the set of techniques "focused on exposing complex AI models to humans in a systematic and interpretable manner" [Samek et al., 2019]. A large body of work on XAI has emerged in recent years [Guidotti et al., 2018b; Bodria et al., 2021]. Counterfactual explanations are used to explain predictions of individual instances in the form: "If X had been different, Y would not have occurred" [Stepin et al., 2021; Karimi et al., 2020a; Schut et al., 2021]. Counterfactual explanations are based on counterfactual examples: modified versions of the input sample that result in an alternative output (i.e., prediction). If the proposed modifications are also actionable, this is referred to as achieving recourse [Ustun et al., 2019; Karimi et al., 2020b].

To motivate our problem, consider an ML application for computational biology: drug discovery is a task that involves generating new molecules that can be used for medicinal purposes [Stokes et al., 2020; Xie et al., 2021]. Given a candidate molecule, a Graph Neural Network (GNN) can predict if this molecule has a certain property that would make it effective in treating a particular disease [Wieder et al., 2020; Guo et al., 2021; Nguyen et al., 2020]. If the GNN predicts it does not have this desirable property, counterfactual explanations can help identify the minimal change required such that the molecule is predicted to have this property. This could help not only inform the design of a new molecule that has this property, but also understand the molecular structures that contribute to this property.

Although GNNs have shown state-of-the-art results on tasks involving graph data [Zitnik et al., 2018; Deac et al., 2019], existing methods for explaining the predictions of GNNs have primarily focused on generating subgraphs that are relevant for a particular prediction [Yuan et al., 2020b; Baldassarre and Azizpour, 2019; Duval and Malliaros, 2021; Lin et al., 2021; Luo et al., 2020; Pope et al., 2019; Schlichtkrull et al., 2021; Vu and Thai, 2020; Ying et al., 2019; Yuan et al., 2021]. However, none of these methods are able to identify the minimal subgraph automatically – they all require the user to specify the size of the subgraph, S, in advance. We show that even if we adapt existing methods to the counterfactual explanation problem, and try varying values for S, such methods are not able to produce valid, accurate counterfactual explanations, and are therefore not well-suited to solve the counterfactual explanation problem. To address this gap, we propose CF-GNNExplainer, a method for generating counterfactual explanations for GNNs.

Similar to other counterfactual methods for tabular or image data proposed in the literature [Verma et al., 2020; Karimi et al., 2020b], CF-GNNExplainer works by perturbing input data at the instance-level. Unlike previous methods, CF-GNNExplainer can generate counterfactual explanations for graph data. In particular, our method iteratively removes edges from the original adjacency matrix based on matrix sparsification techniques, keeping track of the perturbation that leads to a change in prediction, and returning the perturbation with the smallest change w.r.t. the number of edges.

We evaluate CF-GNNExplainer on three public datasets for GNN explanations and measure its effectiveness using four metrics: fidelity, explanation size, sparsity, and accuracy. We find that CF-GNNExplainer is able to generate counterfactual examples with at least 94% accuracy, while removing fewer than 3 edges on average. We make the following contributions:

(1) We formalize the problem of generating counterfactual explanations for GNNs

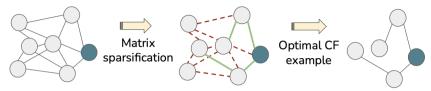


Figure 3.1: Intuition of counterfactual example generation by CF-GNNExplainer.

(Section 3.4).

- (2) We propose CF-GNNExplainer, a novel method for explaining predictions from GNNs (Section 3.5).
- (3) We propose an experimental setup for holistically evaluating counterfactual explanations for GNNs (Section 3.6).

3.2 RELATED WORK

Since our work is a counterfactual XAI approach for GNNs, it is related to GNN explainability (Section 3.2.1) as well as counterfactual explanations (Section 3.2.2). It is also related to adversarial attack methods (Section 3.2.3).

3.2.1 GNN Explainability

Several GNN XAI approaches have been proposed – a recent survey of the most relevant work is presented by Yuan et al. [2020b]. However, unlike our work, *none* of the methods in this survey generate counterfactual explanations.

The majority of existing GNN XAI methods provide an explanation in the form of a subgraph of the original graph that is deemed to be important for the prediction [Yuan et al., 2020b; Baldassarre and Azizpour, 2019; Duval and Malliaros, 2021; Lin et al., 2021; Luo et al., 2020; Pope et al., 2019; Schlichtkrull et al., 2021; Vu and Thai, 2020; Ying et al., 2019; Yuan et al., 2021]. We refer to these as subgraph-generating methods. Such methods are analogous to popular XAI methods such as LIME [Ribeiro et al., 2016a] or SHAP [Lundberg and Lee, 2017], which identify relevant features for a particular prediction for tabular, image, or text data. All of these methods require the user to specify the size of the explanation, S, in advance: the number of features (or edges) to keep. In contrast, CF-GNNExplainer generates counterfactual explanations, which can find the size of the explanation without requiring input from the user. Although both types of techniques are meant for explaining GNN predictions, they are solving fundamentally different problems: counterfactual explanations generate the minimal perturbation such that the prediction changes, while subgraph-retrieving methods identify a relevant (and not necessarily minimal) subgraph that matches the original prediction.

Kang et al. [2019] also generate counterfactual examples for GNNs, but their work focuses on a different task: link prediction. Other GNN XAI methods identify important node features [Huang et al., 2020] or similar examples [Faber et al., 2020]. Yuan et al.

[2020a] and Schnake et al. [2020] generate model-level (i.e., global) explanations for GNNs, which differs from our work since we produce instance-level (i.e., local) explanations.

3.2.2 Counterfactual Explanations

There exists a substantial body of work on counterfactual explanations for tabular, image, and text data [Verma et al., 2020; Karimi et al., 2020b; Stepin et al., 2021]. Some methods treat the underlying classification model as a black-box [Laugel et al., 2018; Guidotti et al., 2018a; Lucic et al., 2020], whereas others make use of the model's inner workings [Tolomei et al., 2017; Wachter et al., 2018; Ustun et al., 2019; Kanamori et al., 2020; Lucic et al., 2022b]. All of these methods are based on perturbing feature values to generate counterfactual examples – they are not equipped to handle graph data with relationships (i.e., edges) between instances (i.e., nodes). In contrast, CF-GNNExplainer provides counterfactual examples specifically for graph data.

3.2.3 Adversarial Attacks

Counterfactual examples are also related to adversarial attacks [Sun et al., 2018]: they both represent instances obtained from minimal perturbations to the input, which induce changes in the prediction made by the learned model. One difference between the two is in the intent: adversarial examples are meant to fool the model, while counterfactual examples are meant to explain the prediction [Freiesleben, 2021; Lucic et al., 2022b]. In the context of graph data, adversarial attack methods typically make minimal perturbations to the *overall graph* with the intention of degrading overall model performance, as opposed to attacking individual nodes. In contrast, we are interested in generating counterfactual examples for individual nodes, as opposed to identifying perturbations to the overall graph. We confirm that the counterfactual examples produced by CF-GNNExplainer are informative and not adversarial by measuring the accuracy of our method (see Section 3.6.3).

3.3 BACKGROUND

In this section, we provide background information on GNNs (Section 3.3.1) and matrix sparsification (Section 3.3.2), both of which are necessary for understanding CF-GNNExplainer.

3.3.1 Graph Neural Networks

Graphs are structures that represent a set of entities (nodes) and their relations (edges). GNNs operate on graphs to produce representations that can be used in downstream tasks such as graph or node classification. The latter is the focus of this work. We refer to the survey papers by Battaglia et al. [2018] and Chami et al. [2021] for an overview of existing GNN methods.

Let $f(A, X; W) \to y$ be any GNN, where y is the set of possible predicted classes, A is an $n \times n$ adjacency matrix, X is an $n \times p$ feature matrix, and W is the learned

weight matrix of f. In other words, A and X are the inputs of f, and f is parameterized by W.

A node's representation is learned by iteratively updating the node's features based on its neighbors' features. The number of layers in f determines which neighbors are included: if there are ℓ layers, then the node's final representation only includes neighbors that are at most ℓ hops away from that node in the graph $\mathcal G$. The rest of the nodes in $\mathcal G$ are not relevant for the computation of the node's final representation. We define the subgraph neighborhood of a node v as the set of the nodes and edges relevant for the computation of f(v) (i.e., those in the ℓ -hop neighborhood of f), represented as a tuple: $\mathcal G_v = (A_v, X_v)$, where A_v is the subgraph adjacency matrix and X_v is the node feature matrix for nodes that are at most ℓ hops away from v. We then define a node v as a tuple of the form $v = (A_v, x)$, where x is the feature vector for v.

3.3.2 Matrix Sparsification

CF-GNNExplainer uses matrix sparsification to generate counterfactual examples, inspired by Srinivas et al. [2017], who propose a method for training sparse neural networks. Given a weight matrix W, a binary sparsification matrix is learned which is multiplied element-wise with W such that some of the entries in W are zeroed out. In the work by Srinivas et al. [2017], the objective is to remove entries in the weight matrix in order to reduce the number of parameters in the model. In our case, we want to zero out entries in the adjacency matrix (i.e., remove edges) in order to generate counterfactual explanations for GNNs. That is, we want to remove the important edges – those that are crucial for the prediction.

3.4 PROBLEM FORMULATION

In general, a counterfactual example \bar{x} for an instance x according to a trained classifier f is found by perturbing the features of x such that $f(x) \neq f(\bar{x})$ [Wachter et al., 2018]. An optimal counterfactual example \bar{x}^* is one that minimizes the distance between the original instance and the counterfactual example, according to some distance function d, and the resulting optimal counterfactual explanation is $\Delta_x^* = \bar{x}^* - x$ [Lucic et al., 2022b].

For graph data, it may not be enough to simply perturb node features, especially since they are not always available. This is why we are interested in generating counterfactual examples by perturbing the graph structure instead. In other words, we want to change the relationships between instances (i.e., nodes), rather than change the instances themselves. Therefore, a counterfactual example for graph data has the form $\bar{v}=(\bar{A}_v,x)$, where x is the feature vector and \bar{A}_v is a perturbed version of A_v , the adjacency matrix of the subgraph neighborhood of a node v. \bar{A}_v is obtained by removing some edges from A_v , such that $f(v) \neq f(\bar{v})$. Following Wachter et al. [2018] and Lucic et al. [2022b], we generate counterfactual examples by minimizing a loss function of the form:

$$\mathcal{L} = \mathcal{L}_{pred}(v, \bar{v} \mid f, g) + \beta \mathcal{L}_{dist}(v, \bar{v} \mid d), \tag{3.1}$$

where v is the original node, f is the original model, g is the counterfactual model that generates \bar{v} , and \mathcal{L}_{pred} is a prediction loss that encourages $f(v) \neq f(\bar{v})$. \mathcal{L}_{dist} is a distance loss that encourages \bar{v} to be close to v, and β controls how important \mathcal{L}_{dist} is compared to \mathcal{L}_{pred} . We want to find \bar{v}^* that minimizes Equation 3.1: this is the optimal counterfactual example for v.

3.5 METHOD: CF-GNNEXPLAINER

To solve the problem defined in Section 3.4, we propose CF-GNNExplainer, which generates $\bar{v}=(\bar{A}_v,x)$ given a node $v=(A_v,x)$. Our method can operate on any GNN model f. To illustrate our method and avoid cluttered notation, let f be a standard, one-layer Graph Convolutional Network [Kipf and Welling, 2017] for node classification:

$$f(A, X; W) = \operatorname{softmax} \left[\tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2} X W \right],$$
 (3.2)

where $\tilde{A} = A + I$, I is the identity matrix, $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$ are entries in the degree matrix \tilde{D} , X is the node feature matrix, and W is the weight matrix [Kipf and Welling, 2017].

3.5.1 Adjacency Matrix Perturbation

First, we define $\bar{A}_v = P \odot A_v$, where P is a binary perturbation matrix that sparsifies A_v . Our aim is to find P for a given node v such that $f(A_v,x) \neq f(P \odot A_v,x)$. To find P, we build upon the method by Srinivas et al. [2017] for training sparse neural networks (see Section 3.3.2), where our objective is to zero out entries in the adjacency matrix (i.e., remove edges). That is, we want to find P that minimally perturbs A_v , and use it to compute $\bar{A}_v = P \odot A_v$. If an element $P_{i,j} = 0$, this results in the deletion of the edge between node i and node j. When P is a matrix of ones, this indicates that all edges in A_v are used in the forward pass.

Similar to the work by Srinivas et al. [2017], we first generate an intermediate, real-valued matrix \hat{P} with entries in [0,1], apply a sigmoid transformation, then threshold the entries to arrive at a binary P: entries greater than or equal to 0.5 become 1, while those below 0.5 become 0. In the case of undirected graphs (i.e., those with symmetric adjacency matrices), we first generate a perturbation vector, which we then use to populate \hat{P} in a symmetric manner, instead of generating \hat{P} directly.

3.5.2 Counterfactual Generating Model

We want our perturbation matrix P to only act on A_v , not \tilde{A}_v , in order to preserve self-loops in the message passing of f. This is because we always want a node representation update to include its own representation from the previous layer. Therefore we first rewrite Equation 3.2 for our illustrative one-layer case to isolate A_v :

$$f(A_v, X_v; W) = \operatorname{softmax} \left[(D_v + I)^{-1/2} (A_v + I) (D_v + I)^{-1/2} X_v W \right].$$
 (3.3)

To generate CFs, we propose a new function g, which is based on f, but it is parameterized by P instead of W. We update the degree matrix D_v based on $P \odot A_v$, add the identity matrix to account for self-loops (as in $\tilde{D_v}$ in Equation 3.2), and call this $\bar{D_v}$:

$$g(A_v, X_v, W; P) = \text{softmax} \left[\bar{D_v}^{-1/2} (P \odot A_v + I) \bar{D_v}^{-1/2} X_v W \right].$$
 (3.4)

In other words, f learns the weight matrix while holding the data constant, while g generates new data points (i.e., counterfactual examples) while holding the weight matrix (i.e., model) constant. Another distinction between f and g is that the aim of f is to find the optimal set of weights that generalizes well on an unseen test set, while the objective of g is to generate an optimal counterfactual example, given a particular node (i.e., \bar{v} is the output of g).

3.5.3 Loss Function Optimization

We generate P by minimizing Equation 3.1, adopting the negative log-likelihood (NLL) loss for \mathcal{L}_{vred} :

$$\mathcal{L}_{pred}(v, \bar{v}|f, g) = \mathbb{1}\left[f(v) = f(\bar{v})\right] \cdot \mathcal{L}_{NLL}(f(v), g(\bar{v})). \tag{3.5}$$

Since we do not want $f(\bar{v})$ to match f(v), we put a negative sign in front of \mathcal{L}_{pred} and include an indicator function to ensure the loss is active as long as $f(\bar{v}) = f(v)$. Note that f and g have the same weight matrix W – the main difference is that g also includes the perturbation matrix P.

 \mathcal{L}_{dist} can be based on any differentiable distance function. In our case, we take d to be the element-wise difference between v and \bar{v} , corresponding to the difference between A_v and \bar{A}_v : the number of edges removed. For undirected graphs, we divide this value by 2 to account for the symmetry in the adjacency matrices. When updating P, we take the gradient of Equation 3.1 with respect to the intermediate \hat{P} , not the binary P.

3.5.4 CF-GNNExplainer

We call our method CF-GNNExplainer and summarize its details in Algorithm 3.1. Given an node in the test set v, we first obtain its original prediction from f and initialize \hat{P} as a matrix of ones, J_n , to initially retain all edges. Next, we run CF-GNNExplainer for K iterations. To find a counterfactual example, we use Equation 3.4.

First, we compute P by thresholding \hat{P} (see Section 3.5.1). Then we use P to obtain the sparsified adjacency matrix that gives us a candidate counterfactual example, \bar{v}_{cand} . This example is then fed to the original GNN, f, and if f predicts a different output than for the original node, we have found a valid counterfactual example, \bar{v} . We keep track of the "best" counterfactual example (i.e., the most minimal according to d), and return this as the optimal counterfactual example \bar{v} after K iterations. Between iterations, we compute the loss following Equations 3.1 and 3.5, and update \hat{P} based on the gradient of the loss. In the end, we retrieve the optimal counterfactual explanation $\Delta_v^* = v - \bar{v}^*$.

Eq 3.1 & 3.5

Undate P

Algorithm 3.1 CF-GNNExplainer: given a node $v=(A_v,x)$ where f(v)=y, generate the minimal perturbation, $\bar{v}=(\bar{A}_v,x)$, such that $f(\bar{v})\neq y$.

Input: node $v=(A_v,x)$, trained GNN model f, CF model g, loss function \mathcal{L} , learning rate α , number of iterations K, distance function d.

```
f(v) = y # Get GNN prediction \hat{P} \leftarrow J_n # Initialization \bar{v}^* = [] for K iterations do \bar{v} = \text{GET\_CF\_EXAMPLE}()
```

```
Function GET_CF_EXAMPLE()
```

 $\mathcal{L} \leftarrow \mathcal{L}(v, \bar{v}, f, g)$

 $\hat{P} \leftarrow \hat{P} + \alpha \nabla_{\hat{P}} \mathcal{L}$

```
\begin{split} P &\leftarrow \operatorname{threshold}(\sigma(\hat{P})) \\ \bar{A}_v &\leftarrow P \odot A_v \\ \bar{v}_{cand} &\leftarrow (\bar{A}_v, x) \\ \text{if } f(v) \neq f(\bar{v}_{cand}) \text{ then } \\ \bar{v} &\leftarrow \bar{v}_{cand} \\ \text{if not } \bar{v}^* \text{ then } \\ \bar{v}^* &\leftarrow \bar{v} & \# \operatorname{First } \operatorname{CF} \\ \text{else if } d(v, \bar{v}) \leq d(v, \bar{v}^*) \text{ then } \\ \bar{v}^* &\leftarrow \bar{v} & \# \operatorname{Best } \operatorname{CF} \\ \text{return } \bar{v}^* \end{split}
```

3.5.5 Complexity

CF-GNN Explainer has time complexity $O(KN^2),$ where N is the number of nodes in the subgraph neighbourhood and K is the number of iterations. We note that high complexity is common for local XAI methods (i.e., SHAP [Lundberg and Lee, 2017], GNN Explainer [Ying et al., 2019], etc.), but in practice, one typically only generates explanations for a subset of the dataset.

3.6 EXPERIMENTAL SETUP

In this section, we outline our experimental setup for evaluating CF-GNNExplainer, including the datasets and models used (Section 3.6.1), the baselines we compare against (Section 3.6.2), the evaluation metrics (Section 3.6.3), and the hyperparameter search method (Section 3.6.4). In total, we run approximately 375 hours of experiments on one Nyidia TitanX Pascal GPU with access to 12GB RAM.

iation.			
	Tree	Tree	BA
	Cycles	Grid	Shapes
# classes	2	2	4
# nodes in motif	6	9	5
# edges in motif (GT)	6	12	6
# nodes in total	871	1231	700
# edges in total	1950	3410	4100
Avg node degree	2.27	2.77	5.87
Avg # nodes in A_v	19.12	30.69	304.40
Avg # edges in A_v	18.99	33.94	1106.24

Table 3.1: Dataset statistics. The # edges in the motif indicates the size of the ground truth (GT) explanation.

3.6.1 Datasets and Models

Given the challenges associated with defining and evaluating the accuracy of XAI methods [Doshi-Velez and Kim, 2018], we first focus on synthetic tasks where we know the ground-truth explanations. Although there exist real graph classification datasets with ground-truth explanations [Debnath et al., 1991], there do not exist any real node classification datasets with ground-truth explanations, which is the task we focus on in this chapter. Building such a dataset would be an excellent contribution, but is outside the scope of this paper.

In our experiments, we use the TREE-CYCLES, TREE-GRIDS, BA-SHAPES datasets from the work by Ying et al. [2019]. These datasets were created specifically for the task of explaining node classification predictions from GNNs. Each dataset consists of (i) a base graph, (ii) motifs that are attached to random nodes of the base graph, and (iii) additional edges that are randomly added to the overall graph. They are all undirected graphs. The classification task is to determine whether or not the nodes are part of the motif. The purpose of these datasets is to have a ground-truth for the "correctness" of an explanation: for nodes in the motifs, the explanation is the motif itself [Luo et al., 2020]. The dataset statistics are available in Table 3.1.

TREE-CYCLES consists of a binary tree base graph with 6-cycle motifs, TREE-GRIDS also has a binary tree as its base graph, but with 3×3 grids as the motifs. For BA-SHAPES, the base graph is a Barabasi-Albert (BA) graph with house-shaped motifs, where each motif consists of 5 nodes (one for the top of the house, two in the middle, and two on the bottom). Here, there are four possible classes (not in motif, in motif: top, middle, bottom). We note that compared to the other two datasets, the BA-SHAPES dataset is much more densely connected – the node degree is more than twice as high as that of the TREE-CYCLES or TREE-GRID datasets, and the average number of nodes and edges in each node's computation graph is order(s) of magnitude larger. We use the same experimental setup (i.e., dataset splits, model architecture) as Ying et al. [2019] to train a 3-layer GCN (hidden size = 20) for each task. Our GCNs have at least 87% accuracy on the test set.

3.6.2 Baselines

Since existing GNN XAI methods give explanations in the form of relevant subgraphs as opposed to counterfactual examples, it is not straightforward to identify baselines for our experiments that ensure a fair comparison between methods. To evaluate CF-GNNExplainer, we compare against 4 baselines: RANDOM, 1HOP, RM-1HOP, and GNNExplainer. The random perturbation is meant as a sanity check. We randomly initialize the entries of $\hat{P} \in [-1,1]$ and apply the same sigmoid transformation and thresholding as described in Section 3.5.1. We repeat this K times and keep track of the most minimal perturbation resulting in a counterfactual example. Next, we compare against baselines that are based on the ego graph of v (i.e., its 1-hop neighbourhood): 1HOP keeps all edges in the ego graph of v, while RM-1HOP removes all edges in the ego graph of v.

Our fourth baseline is based on GNNExplainer by Ying et al. [2019], which identifies the S most relevant edges for the prediction (i.e., the most relevant subgraph of size S). To generate counterfactual explanations, we remove the subgraph generated by GNNExplainer. We include this method in our experiments in order to have a baseline based on a prominent GNN XAI method, but we note that subgraph-retrieving methods like GNNExplainer are not meant for generating counterfactual explanations. Unlike our method, GNNExplainer cannot automatically find a *minimal* subgraph and therefore requires the user to determine the number of edges to keep in advance (i.e., the value of S). As a result, we cannot evaluate how minimal its counterfactual explanations are, but we can compare it against our method in terms of (i) its ability to generate valid counterfactual examples, and (ii) how accurate those counterfactual examples are. We report results on GNNExplainer for $S \in \{1, 2, 3, 4, 5, GT\}$, where GT is the size of the ground truth explanation (i.e., the number of edges in the motif, see Table 3.1).

3.6.3 Metrics

We generate a counterfactual example for each node in the graph separately and evaluate in terms of four metrics.

Fidelity: is defined as the proportion of nodes where the original predictions match the prediction for the explanations [Molnar, 2019; Ribeiro et al., 2016a]. Since we generate counterfactual examples, we do not want the original prediction to match the prediction for the explanation, so we want a low value for fidelity.

Explanation Size: is the number of removed edges. It corresponds to the \mathcal{L}_{dist} term in Equation 3.1: the difference between the original A_v and the counterfactual \bar{A}_v . Since we want to have *minimal* explanations, we want a small value for this metric. Note that we cannot evaluate this metric for GNNEXPLAINER.

Sparsity: measures the proportion of edges in A_v that are removed [Yuan et al., 2020b]. A value of 0 indicates all edges in A_v were removed. Since we want *minimal* explanations, we want a value close to 1. Note that we cannot evaluate this metric for GNNEXPLAINER.

Accuracy: is the mean proportion of explanations that are "correct". Following Ying et al. [2019]; Luo et al. [2020], we only compute accuracy for nodes that are originally predicted as being part of the motifs, since accuracy can only be computed on instances for which we know the ground truth explanations. Given that we want *minimal* explanations, we consider an explanation to be correct if it *exclusively* involves edges that are inside the motifs (i.e., only removes edges that are within the motifs).

3.6.4 Hyperparameter Search

We experiment with different optimizers and hyperparameter values for the number of iterations K, the trade-off parameter β , the learning rate α , and the Nesterov momentum m (when applicable). We choose the setting that produces the most counterfactual examples. We test the number of iterations $K \in \{100, 300, 500\}$, the trade-off parameter $\beta \in \{0.1, 0.5\}$, the learning rate $\alpha \in \{0.005, 0.01, 0.1, 1\}$, and the Nesterov momentum $m \in \{0, 0.5, 0.7, 0.9\}$. We test Adam, SGD and AdaDelta as optimizers. We find that for all three datasets, the SGD optimizer gives the best results, with k = 500, $\beta = 0.5$, and $\alpha = 0.1$. For the tree-cycles and tree-grid datasets, we set m = 0, while for the ba-shapes dataset, we use m = 0.9.

3.7 RESULTS

We evaluate CF-GNNExplainer in terms of the metrics outlined in Section 3.6.3. The results are shown in Table 3.2 and Table 3.3. In cases where the baselines outperform CF-GNNExplainer on a particular metric, they perform poorly on the rest of the metrics, or on other datasets.

3.7.1 Main Findings

Fidelity: CF-GNNExplainer outperforms 1HOP across all three datasets, and outperforms RM-1HOP for TREE-CYCLES and TREE-GRID in terms of fidelity. We find that RANDOM has the lowest fidelity in all cases – it is able to find counterfactual examples for every single node. In the following subsections, we will see that this corresponds to poor performance on the other metrics.

Table 3.2: Results comparing our method (abbreviated as CF-GNN) to RANDOM, 1HOP, and RM-1HOP. Below each metric, ▼ indicates a low value is desirable, while ▲ indicates a high value is desirable.

TREE-CYCLES				TREE	-GRID	BA-SHAPES						
Method	Fid. ▼	Size ▼	Spar.	Acc.	Fid. ▼	Size ▼	Spar.	Acc.	Fid. ▼	Size ▼	Spar.	Acc.
RANDOM	0.00	4.70	0.79	0.63	0.00	9.06	0.75	0.77	0.00	503.31	0.58	0.17
1нор км-1нор	0.32 0.46	15.64 2.11	0.13 0.89	0.45 —	0.32 0.61	29.30 2.27	0.09 0.92	0.72 —	0.60 0.21	504.18 10.56	0.05 0.97	0.18 0.99
CF-GNN	0.21	2.09	0.90	0.94	0.07	1.47	0.94	0.96	0.39	2.39	0.99	0.96

Explanation Size: Figures 3.2 to 3.5 show histograms of the explanation size for CF-GNNExplainer and the baselines. We see that across all three datasets, CF-GNNExplainer has the smallest (i.e., most minimal) explanation sizes. This is especially true when comparing to random and 1hop for the ba-shapes dataset, where we had to use a different scale for the x-axis due to how different the explanation sizes were. We postulate that this difference could be because ba-shapes is a much more densely connected graph; it has fewer nodes but more edges compared to the other two datasets, and the average number of nodes and edges in the subgraph neighborhood is order(s) of magnitude larger (see Table 3.1). Therefore, when performing random perturbations, there is substantial opportunity to remove edges that do not necessarily need to be removed, leading to much larger explanation sizes. When there are many edges in the subgraph neighborhood, removing everything except the 1-hop neighbourhood, as is done in 1hop, also results in large explanation sizes. In contrast, the loss function used by CF-GNNExplainer ensures that only a few edges are removed, which is the desirable behavior since we want minimal explanations.

Sparsity: CF-GNNExplainer outperforms the RANDOM, RM-1HOP, 1HOP baselines for all three datasets in terms of sparsity. We note CF-GNNExplainer and RM-1HOP perform much better on this metric in comparison to the other methods, which aligns with the results from explanation size.

Accuracy: We observe that CF-GNNExplainer has the highest accuracy for the TREE-CYCLES and TREE-GRID datasets, whereas RM-1HOP has the highest accuracy for BA-SHAPES. However, we are unable to calculate the accuracy of RM-1HOP for the other two datasets since it is unable to generate *any* counterfactual examples for motif nodes, contributing to the low sparsity on those datasets. We observe accuracy levels upwards of 94% for CF-GNNExplainer across *all* datasets, indicating that it is consistent in correctly removing edges that are crucial for the initial predictions in the vast majority of cases (see Table 3.2).

3.7.2 Comparison to GNNEXPLAINER

Table 3.3 shows the results comparing our method to GNNExplainer. We find that our method outperforms GNNExplainer across all three datasets in terms of both fidelity and accuracy, for all tested values of S. However, this is not surprising since GNNExplainer is not meant for generating counterfactual explanations, so we cannot expect it to perform well on a task it was not designed for. We cannot compare explanation size or sparsity fairly since GNNExplainer requires the user to input the value of S.

3.7.3 Summary of results

Evaluating on four distinct metrics for each dataset gives us a more holistic view of the results. We find that across all three datasets, CF-GNNExplainer can generate counterfactual examples for the majority of nodes in the test set (i.e., low fidelity), while only removing a small number of edges (i.e., low explanation size, high sparsity).

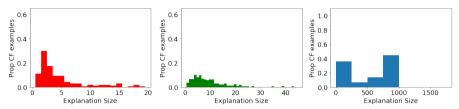


Figure 3.2: Histograms showing the proportion of counterfactual examples that have a certain explanation size from random. Note the x-axis for BA-SHAPES goes up to 1500. Left: TREE-CYCLES, Middle: TREE-GRID, Right: BA-SHAPES.

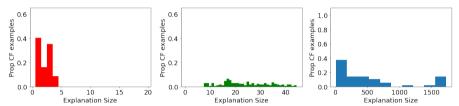


Figure 3.3: Histograms showing the proportion of counterfactual examples that have a certain explanation size from 1hop. Note the x-axis for BA-shapes goes up to 1500. Left: Tree-cycles, Middle: Tree-grid, Right: BA-shapes.

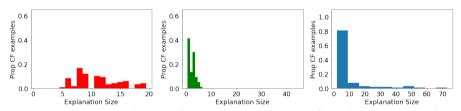


Figure 3.4: Histograms showing the proportion of counterfactual examples that have a certain explanation size from RM-1HOP. Note the x-axis for BA-SHAPES goes up to 70. Left: TREE-CYCLES, Middle: TREE-GRID, Right: BA-SHAPES.

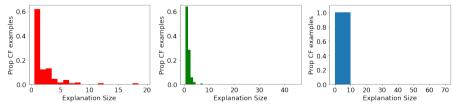


Figure 3.5: Histograms showing the proportion of counterfactual examples that have a certain explanation size from CF-GNNExplainer. Note the *x*-axis for BA-SHAPES goes up to 70. Left: TREE-CYCLES, Middle: TREE-GRID, Right: BA-SHAPES.

Table 3.3: Results comparing our method to GNNEXPLAINER. GNNEXPLAINER cannot find S automatically, so we try varying values of S. GT indicates the size of the ground truth explanation for each dataset. CF-GNNExplainer finds S automatically. Below each metric, \blacksquare indicates a low value is desirable, while \blacksquare indicates a high value is desirable.

	TREE-CYCLES				TREE	-GRID		BA-SHAPES				
GNNExp	Fid. ▼	Size ▼	Spars.	Acc.	Fid. ▼	Size ▼	Spars.	Acc.	Fid. ▼	Size ▼	Spars.	Acc. ▲
S = 1	0.65	1.00	0.92	0.61	0.69	1.00	0.96	0.79	0.90	1.00	0.94	0.52
S = 2	0.59	2.00	0.85	0.54	0.51	2.00	0.92	0.78	0.85	2.00	0.91	0.40
S = 3	0.56	3.00	0.79	0.51	0.46	3.00	0.88	0.79	0.83	3.00	0.87	0.34
S = 4	0.58	4.00	0.72	0.48	0.42	4.00	0.84	0.79	0.83	4.00	0.83	0.31
S = 5	0.57	5.00	0.66	0.46	0.40	5.00	0.80	0.79	0.81	5.00	0.81	0.27
S = GT	0.55	6.00	0.57	0.46	0.35	11.83	0.53	0.74	0.82	6.00	0.79	0.24
CF-GNN	0.21	2.09	0.90	0.94	0.07	1.47	0.94	0.96	0.39	2.39	0.99	0.96

For nodes where we know the ground truth (i.e., those in the motifs) we achieve at least 94% accuracy.

Although random can generate counterfactual examples for every node, they are not very minimal or accurate. The latter is also true for 1HOP – in general, it has the worst scores for explanation size, sparsity and accuracy. GNNExplainer performs at a similar level as 1HOP, indicating that although it is a prominent GNN XAI method, it is not well-suited for solving the counterfactual explanation problem.

RM-1HOP is competitive in terms of explanation size, but it performs poorly in terms of fidelity for the TREE-CYCLES and TREE-GRID datasets, and its accuracy on these datasets is unknown since it is unable to generate *any* counterfactual examples for nodes in the motifs. These results show that our method is simple and extremely effective in solving the counterfactual explanation task, unlike the baselines we test.

3.8 SOCIETAL IMPACT

Researchers have raised concerns about the hidden assumptions behind the use of counterfactual examples [Barocas et al., 2020], as well as potentials for misuse [Kasirzadeh and Smart, 2021]. When explaining ML systems through counterfactual examples, it is crucial to account for the context in which the systems are deployed. Counterfactual explanations are not a guarantee to achieving recourse [Ustun et al., 2019] – changes suggested should be seen as candidate changes, not absolute solutions, since what is pragmatically actionable differs depending on the context.

We believe it is crucial for the ML community to invest in developing more rigorous evaluation protocols for XAI methods. We suggest that researchers in XAI collaborate with researchers in human-computer interaction to design human-centered user studies about evaluating the utility of XAI methods in practice. We are glad to see initiatives for such collaborations already taking place [Ehsan et al., 2021].

3.9 CONCLUSION

In this chapter, we propose CF-GNNExplainer, a method for generating counterfactual explanations for any GNN. Our simple and effective method is able to generate counterfactual explanations that are (i) minimal, both in terms of the absolute number of edges removed (explanation size), as well as the proportion of the subgraph neighborhood that is perturbed (sparsity), and (ii) accurate, in terms of removing edges that we know to be crucial for the initial predictions.

We evaluate our method on three commonly used datasets for GNN explanation tasks and find that these results hold across all three datasets. We find that existing GNN XAI methods are not well-suited to solving the counterfactual explanation task, while CF-GNNExplainer is able to reliably produce minimal, accurate counterfactual explanations.

Should we switch order of last two paragraphs here?

In its current form, CF-GNNExplainer is limited to performing edge deletions in the context of node classification tasks. For future work, we plan to incorporate node feature perturbations in our framework and extend CF-GNNExplainer to accommodate graph classification tasks. We also plan to investigate adapting graph attack methods for generating counterfactual explanations, as well as conduct a user study to determine if humans find CF-GNNExplainer useful in practice.

This answers **RQ2**: we can generate counterfactual explanations for graph-based models by extending the problem formalization from Chapter 2 to accommodate graph data, then applying similar gradient-based optimization techniques for each example in the dataset. This can be used to explain predictions from any GNN. In the following section, we will investigate how to generate and evaluate explanations in a human-centric manner.

REPRODUCIBILITY

To facilitate the reproducibility of the work in this chapter, our code is available at https://github.com/a-lucic/cf-gnnexplainer.

Part II

Humans

Contrastive Explanations for Retail Forecasting

In this part of the thesis, we shift our focus to the consumers of explanations: humans. In this chapter, we address the following research question:

RQ3: How can we develop an explanation method based on a real-world use case and evaluate it in a human-centric way?

We answer **RQ3** by first identifying a use case where users seek explanations: understanding errors in sales forecasting. We then design an algorithm that generates explanations for large errors in regression predictions based on Monte Carlo simulations. To evaluate our method, we design a user study that includes both objective and subjective components, where we contrast and compare the results between two types of users: researchers and practitioners. Our experimental results show that explanations generated by our method help both types of users understand why large errors in predictions occur, but do not have an impact on their trust or confidence in the model.

4.1 INTRODUCTION

As more and more decisions about humans are made by machines, it becomes imperative to understand how these outputs are produced and what drives a model to a particular prediction [Ribeiro et al., 2016b]. As a result, algorithmic interpretability has gained significant interest and traction in the ML community over the past few years [Doshi-Velez and Kim, 2018]. However, there exists considerable skepticism outside of the ML community due to a perceived lack of transparency behind algorithmic predictions, especially when errors are produced [Dietvorst et al., 2015]. We aim to evaluate the effect of explaining model outputs, specifically large errors, on users' attitudes towards trusting and deploying complex, automatically learned models.

Further motivation for interpretable ML is provided by significant societal developments. Important examples include the recently enacted European General Data

This chapter was published at the ACM Conference on Fairness, Accountability, and Transparency (FAccT 2020) under the title "Why Does My Model Fail: Contrastive Explanations for Retail Forecasting" [Lucic et al., 2020], where it won a best paper award.

Protection Regulation (GDPR), which specifies that individuals will have the right to "the logic involved in any automatic personal data processing" [EU, 2016]. In Canada and the United States, this right to an explanation is an integral part of financial regulations, which is why banks have not been able to use high-performing "black-box" models to evaluate the credit-worthiness of their customers. Instead, they have been confined to easily interpretable algorithms such as decision trees (for segmenting populations) and logistic regression (for building risk scorecards) [Khandani et al., 2010]. At NeurIPS 2017, an Explainable ML Challenge was launched to combat this limitation, indicating the finance industry's interest in exploring algorithmic explanations [FICO, 2017].

We use explanations as a mechanism for supporting innovation and technological development while keeping the human "in the loop" by focusing on predictive modeling as a tool that aids individuals with a given task. Specifically, our interest lies with interpretability in a scenario where users with varying degrees of ML expertise are confronted with large errors in the outcome of predictive models. We focus on explaining large errors because people tend to be more curious about unexpected outcomes rather than ones that confirm their prior beliefs [Hilton and Slugoski, 1986].

However, Dietvorst et al. [2015] showed that when users are confronted with errors in algorithmic predictions, they are less likely to use the model. Seeing an algorithm make mistakes significantly decreases confidence in the model, and users are more likely to choose a human forecaster instead, even after seeing the algorithm outperform the human [Dietvorst et al., 2015]. This indicates that prediction mistakes have a significant impact on users' perception of the model. By focusing on explaining mistakes, we hope to give insight into this phenomenon of algorithm aversion while also giving users the types of explanations they are interested in seeing.

Our work was motivated by the needs of analysts at a large retailer in the Netherlands, working on sales forecasting. Current models in production are based on simple autoregressive methods, but there is an interest in exploring more complex techniques. However, the added complexity comes at the expense of interpretability, which is problematic for the company, especially when a complex model produces a forecast that is very different from the actual target value. This leads us to focus on explaining errors in regression predictions in this work. However, it should be noted that our method can be extended to classification predictions by defining "distances" between classes or by simply defining all errors as large errors.

We focus on two aspects of explainability in this scenario: the *generation* of explanations of large errors and the corresponding *effectiveness* of these explanations. Prior methods for generating explanations fail at generating explanations for large errors because they produce similar explanations for predictions resulting in large errors and those resulting in reasonable predictions (see Table 4.2 in Section 4.4 for an example). We propose a method for explaining large prediction errors, called Monte Carlo Bounds for Reasonable Predictions (MC-BRP), that shows users:

- (i) The required bounds of the most important features in order to have a prediction resulting in a reasonable prediction.
- (ii) The relationship between each of these features and the target.

It should be noted that in this chapter, we focus on explaining errors *in hindsight*, that is, we examine large errors once they have occurred and are not predicting them in advance without having access to the ground truth. We are also not using these explanations to improve the model, but rather examine the effectiveness of explaining large errors via MC-BRP on users' trust in the model and attitudes towards deploying it, as well as their understanding of the explanations. We test on a wide range of users, including both Practitioners and Researchers, and analyze the differences in attitudes between these users. We also reflect on the process of conducting a user study by outlining some limitations of our study and make some recommendations for future work.

We address the following research questions:

RQ3.1: Are the contrastive explanations generated by MC-BRP about large errors in predictions (i) interpretable, or (ii) actionable? More specifically,

- (i) Can contrastive explanations about large errors give users enough information to simulate the model's output (forward simulation)?
- (ii) Can such explanations help users understand the model such that they can manipulate an observation's input values in order to change the output (counterfactual simulation)?

RQ3.2: How does providing contrastive explanations generated by MC-BRP for large errors impact users' perception of the model? Specifically, we want to investigate the following:

- (i) Does being provided with contrastive explanations generated by MC-BRP impact users' understanding of why the model produces errors?
- (ii) Does it impact their willingness to deploy the model?
- (iii) Does it impact their level of trust in the model?
- (iv) Does it impact their confidence in the model's performance?

Consequently, we make the following contributions:

- We contribute a method, MC-BRP, for generating contrastive explanations specifically for large errors in regression tasks.
- We evaluate our explanations through a user study with 75 participants in both objective and subjective terms.
- We conduct an analysis on the differences in attitudes between Practitioners and Researchers.

In Section 4.2 we discuss related work and identify how our problem relates to the current literature. In Section 4.3 we formally describe the methodology of explanations based on MC-BRP and in Section 4.4 we motivate our choice of dataset and describe the user study setup. In Section 4.5 we detail the results of the user study; we conduct further analyses in Section 4.6. In Section 4.7 we conclude and make recommendations for future work.

4.2 RELATED WORK

Guidotti et al. [2018b] compile a survey of current methods in interpretable machine learning and develop a taxonomy for classifying methods using four criteria:

· Problem:

- (i) Model explanations: interpret black-box model as a whole (globally)
- (ii) Outcome explanations: interpret individual black-box predictions (locally)
- (iii) *Inspection*: interpret model behavior through visual representations (globally or locally)
- (iv) Transparent design: model is inherently interpretable (globally or locally)
- Model: neural networks, tree ensembles, SVMs, model-agnostic
- Explanator: decision trees/rules, feature importances, salient masks, sensitivity analysis, partial dependence plots, prototype selection, neuron activation
- · Data: tabular, image or text

Based on this schema, our setting is this chapter is an *outcome explanation* problem for *tree ensembles*. We use *sensitivity analysis*, specifically Monte Carlo simulations, on *tabular* data to generate our explanations.

Existing work on generating outcome explanations specifically for tree ensembles involves finding counterfactual examples [Tolomei et al., 2017], identifying influential training samples [Sharchilev et al., 2018], or identifying important features [Lundberg et al., 2020]. Importantly, none of these publications are specifically about (i) explaining errors, or (ii) explaining regressions. On the contrary, these publications are all based on binary classification tasks and the explanations do not necessarily provide insight into prediction mistakes.

Tolomei et al. [2017] propose a method for generating counterfactual examples by identifying decision paths of interest that would result in a different prediction, then traversing down each of these paths and perturbing the instance x such that it satisfies the path in question. If this perturbation, x', (i) satisfies the decision path, and (ii) changes the prediction in the overall ensemble, then it is a candidate transformation of x. After computing all possible candidate transformations by traversing over all paths of interest (i.e., those leading to a different prediction), the candidate transformation with the smallest distance from x is selected as the counterfactual example. The explanation, then, is the difference between x and x'. Although Tolomei et al. [2017]'s method also produces contrastive explanations, our method differs from theirs since we are not aiming to identify one counterfactual example, but rather a range of feature values for which the prediction would be different. Another difference is that we do not assume full access to the original model.

Sharchilev et al. [2018] also generate outcome explanations for tree ensembles. Their methodology is based on finding influential training samples in order to automatically improve the model, which differs from our work since their explanations are not of a contrastive nature. These influential training samples help us understand

why a certain class was predicted for a given instance, but they make no reference to the alternative class(es). It should be noted that they include a use case on identifying harmful training examples — ones that contributed to incorrect predictions — which can be seen as a way to explain errors.

Lundberg et al. [2020] propose a method for determining how much each feature contributes to a prediction and present a ranked list of the most important features as the explanation. The approach is based on the computationally intensive Shapley values [Lundberg and Lee, 2017], for which the authors develop a tree-specific approximation. This differs from our method since identifying the most important features is only a preliminary step in our pipeline — our work extends beyond this by including (i) feature bounds that result in reasonable predictions, and (ii) the relationship between the features and the target as a tool to help users inspect what goes wrong when the prediction error is large.

Ribeiro et al. [2016a] also propose a method for identifying local feature importances and this is the one we use in our pipeline. Their method, LIME, is model-agnostic and is based on approximating the original model locally with a linear model. We share their objective of evaluating users' attitudes towards a model through local explanations but we further specify our task as explaining instances where there are large errors in predictions. Based on preliminary experiments, we find that LIME is insufficient for our task setting for two reasons:

- (i) For regression tasks, LIME's approximation of the original model is not exact. This "added" error can be quite large given that our target is typically of order 10^6 , and this convolutes our definition of a large error.
- (ii) The features LIME deems most important are similar regardless of whether the prediction results in a large error or not, which does not provide any specific insight into why a large error occurs. These experiments are detailed in Section 4.4.

Other work on contrastive explanations includes identifying features that should be present or absent in order to justify a classification [Dhurandhar et al., 2018; Hendricks et al., 2018] or model-agnostic counterfactuals [Wachter et al., 2018; Russell, 2019]. These all differ from our method since they are not specifically about explaining errors. Furthermore, the work by Dhurandhar et al. [2018] and Hendricks et al. [2018] is based on the binary presence or absence of input features, whereas our method perturbs inputs instead of removing them altogether.

Our work in this chapter can also be viewed as a form of outlier detection. However, it differs from the standard literature outlined by Pimentel et al. [2014] with respect to the objective: we are not necessarily trying to identify outliers in terms of the training data but rather explain instances in the test set whose errors are so large that they are considered to be anomalies.

Miller et al. [2017] perform a survey of the papers cited in the "Related Works" section of the call for the IJCAI 2017 Explainable AI workshop [IJCAI, 2017] and find that the majority do not base their methods on the available research about explanations from other disciplines such as philosophy, psychology or cognitive sciences, or evaluate on real users. In contrast, our method is rooted in the corresponding philosophical

literature [Hilton, 1990; Lipton, 1990; Hilton and Slugoski, 1986] and our evaluation is based on a user study.

4.3 METHOD

The intuition behind MC-BRP is based on identifying the unusual properties of a particular observation. We make the assumption that large errors occur due to unusual feature values in the test set that were not common in the training set.

Given an observation that results in a large error, MC-BRP generates a set of bounds for each feature that would result in a reasonable prediction as opposed to a large error. We also include the trend as part of the explanation in order to help users understand the relationship between each feature and the target, and how the input should be changed in order to change the output.

As pointed out previously, we consider our task of identifying and explaining large errors somewhat similar to that of an outlier detection problem. A standard definition of a statistical outlier is an instance that falls outside of a threshold based on the interquartile range. A widely used version of this, called Tukey's fences, is defined as follows [Tukey, 1977]:

$$[Q_1 - 1.5(Q_3 - Q_1), Q_3 + 1.5(Q_3 - Q_1)],$$

where Q_1 and Q_3 are the first and third quartiles, respectively.

Definition 4.3.1. Let x be an observation in the test set X and let t, \hat{t} be the actual and predicted target values of x, respectively. Let ϵ be the corresponding prediction error for x, and let E be the set of all errors of X. Then ϵ is a *large error* iff

$$\epsilon > Q_3(E) + 1.5(Q_3(E) - Q_1(E)),$$

where $Q_1(E), Q_3(E)$ are the first and third quartiles of the set of errors, respectively. We denote this threshold as ϵ_{large} .

We can view X in Definition 4.3.1 as a disjoint union of two sets:

- (i) R: the set of observations resulting in reasonable predictions, and
- (ii) L: the set of observations resulting in large errors.

We determine the n most important features based on LIME $\Phi^{(x)} = \{\phi_j^{(x)}\}_{j=1}^n$, for all $x \in X$. It should be noted there exist alternative methods for determining the most important features for a particular prediction [Lundberg and Lee, 2017], which would also be appropriate.

Given $x\in X$, for each $\phi_j^{(x)}\in\Phi^{(x)}$, we determine two sets of characteristics through Monte Carlo simulations:

(i)
$$[a_{\phi_j^{(x)}},b_{\phi_j^{(x)}}]$$
: the bounds for values of $\phi_j^{(x)}$ such that $x\in R, x\not\in L$.

Table 4.1: An example of an explanation generated by MC-BRP. Here, each of the input
values is outside of the range required for a reasonable prediction, which explains why
this particular prediction results in a large error.

Input	Definition	Trend	Value	Reasonable range
Α	total_contract_hrs	As input \uparrow , sales \uparrow	9628.0	[4140,6565]
В	advertising_costs	As input ↑, sales ↑	18160.7	[8290,15322]
C	num_transactions	As input ↑, sales ↑	97332.0	[51219,75600]
D	total_headcount	As input ↑, sales ↑	226.0	[95,153]
E	floor_surface	As input \uparrow , sales \uparrow	2013.6	[972,1725]

(ii) $\rho_{\phi_{j}^{(x)}}$: the relationship between $\phi_{j}^{(x)}$ and the target, t.

We perturb the feature values for $l \in L$ using Monte Carlo simulations in order to determine what feature values are required to produce a reasonable prediction. The algorithm for determining R', the set of Monte Carlo simulations resulting in reasonable predictions, is detailed in Algorithm 4.1.

Given $l \in L$, we determine Tukey's fences for each feature in $\Phi^{(l)}$ based on the feature values from R. This gives us the bounds from which we sample for our feature perturbations.

Next, we randomly sample from these bounds for each $\phi_j^{(l)} \in \Phi^{(l)}$ m-times to generate mn versions of our original observation, l. We call the i-th perturbed version l'_i , where $i \in \{1, \ldots, mn\}$.

We then test the original model f on each l'_i , obtain a new prediction, $\hat{t'_i}$, and construct R', the set of perturbations resulting in reasonable predictions.

Once R' is generated, we compute the mean, standard deviation and Pearson coefficient [Swinscow, 1997] of the top n features of $l \in L$, $\Phi^{(l)}$, based on this set.

Algorithm 4.1 Monte Carlo simulation: creates a set of perturbed instances resulting in reasonable predictions R' for each large error $l \in L$

Input: instance l, set of l's most important features $\Phi^{(l)}$, 'black-box' model f, large

error threshold ϵ_{large} , number of MC perturbations per feature m. $R' = \emptyset$ for all $\phi_j^{(l)}$ in $\Phi^{(l)}$ do $TF(\phi_j^{(l)}) \leftarrow \text{Tukeys_fences}(\phi_j^{(l)}) \qquad \text{\#Based on } R$ for i in range (0,m) do $\phi_j'^{(l)} \leftarrow \text{Random_sample}(TF(\phi_j^{(l)}))$ $l_i' \leftarrow l_i.\text{Replace}(\phi_j^{(l)},\phi_j'^{(l)})$ $l_i' \leftarrow f(l_i') \qquad \text{\#New prediction}$ if $|\hat{t}_i' - t_i| < \epsilon_{large}$ then $R' \leftarrow R' \cup l_i'$

return R'

Definition 4.3.2. The *trend*, $\rho_{\phi_j^{(x)}}$, of each feature is the Pearson coefficient between each feature $\phi_j^{(x)}$ and the predictions \hat{t}_i' based on the observations in R'. It is a measure of linear correlation between two variables [Swinscow, 1997].

The set of bounds for each feature in $\Phi^{(x)}$ such that \hat{t} results in a reasonable prediction are based on the mean and standard deviation of each $\phi_i^{(x)} \in \Phi^{(x)}$.

Definition 4.3.3. The *reasonable bounds* for values of each feature ϕ_j in $\Phi^{(x)}$, $[a_{\phi_j^{(x)}}, b_{\phi_j^{(x)}}]$, are

 $\left[\mu(\phi_j^{(x)}) - \sigma(\phi_j^{(x)}), \mu(\phi_j^{(x)}) + \sigma(\phi_j^{(x)})\right],$

where $\mu(\phi_j^{(x)})$ and $\sigma(\phi_j^{(x)})$ are the mean and standard deviation of each feature, respectively, based on R'.

We compute the trend and the reasonable bounds for each of the n most important features and present them to the user in a table. Table 4.1 shows an example of an explanation generated by MC-BRP; the dataset used for this example is detailed in Section 4.4.1.

4.4 EXPERIMENTAL SETUP

Current explanation methods mostly serve individuals with ML expertise [Guidotti et al., 2018b], but they should be extended to cater to users outside of the ML community [Miller, 2019]. Unlike previous work, our method, MC-BRP, generates contrastive explanations by framing the explanation around the prediction error, and aims to help users understand (i) what contributed to the large error, and (ii) what would need to change in order to produce a reasonable prediction. Presenting explanations in a contrastive manner helps frame the problem and narrows the user's focus regarding the possible outcomes [Hilton, 1990; Lipton, 1990].

Our explanations are contrastive because they display to the user what would have needed to change in the input order to obtain an alternative outcome from the model — in other words, why this prediction results in a large error as opposed to a reasonable prediction.

4.4.1 Dataset and Model

Our task is predicting monthly sales of the company's stores with 45 features including financial, workforce and physical store aspects. Since not all of our Practitioners have experience with ML, using an internal dataset with familiar features allows them to leverage some of their domain expertise. The dataset includes 45628 observations from 563 stores, collected at four-week intervals spanning from 2010–2015. We split the data by year (training: 2010–2013, test: 2014–2015) to simulate a production environment, and we treat every unique combination of store, interval and year as an independent observation. After preprocessing, we have 21415 and 12239 observations in our training and test sets, respectively. We train the gradient boosting regressor from scikit-learn with the default settings and obtain an \mathbb{R}^2 of 0.96.

We verify our assumption that large errors are a result of unusual feature values by generating MC-BRP explanations for all instances in our test set using n=5 features and m=10000 Monte Carlo simulations. In our dataset, we find that 48% of instances resulting in large errors have feature values outside the reasonable range for all of the n=5 most important features, compared to only 24% of instances resulting in reasonable predictions. Although this is not perfect, it is clear that MC-BRP produces explanations that are at least somewhat able to distinguish between these two types of predictions.

4.4.2 Comparison to LIME

Hilton [2017] states that explanations are selective – it is not necessary or even useful to state all the possible causes that contributed to an outcome. The significant part of an explanation is what distinguishes it from the alternative outcome. If LIME explanations were suitable for our problem, then we would expect to see different features deemed important for instances resulting in large errors compared to those resulting in acceptable errors. This would help the user understand why a particular prediction resulted in a large error.

However, when generating LIME explanations for our test set using n=5 features, we do not see much of a distinction in the most important features between predictions that result in large errors and those that do not. For example, advertising_costs is one of the top 5 most important features in 18.8% of instances with large errors and 18.7% of instances with reasonable predictions. These results are summarized in Table 4.2.

Table 4.2: The top $n =$	5 features according to LIME for observations resulting in large
errors vs. reasonable	predictions.

Large errors		Reasonable Predictions		
advertising_costs	0.188	advertising_costs	0.187	
total_contract_hrs	0.175	total_contract_hrs	0.179	
$num_transactions$	0.151	$num_transactions$	0.156	
floor_surface	0.124	total_headcount	0.134	
total_headcount	0.123	floor_surface	0.122	
month	0.109	month	0.094	
mean_tenure	0.046	mean_tenure	0.046	
earnings_index	0.033	earnings_index	0.031	

Furthermore, we originally tried to design our control group user study using explanations from LIME, but found that test users from the company could not make sense of the objective questions about prediction errors because LIME does not provide any insight about errors specifically. Given that we could not even ask questions about errors using LIME explanations to users without confusing them, it is clear that LIME is inappropriate for our task.

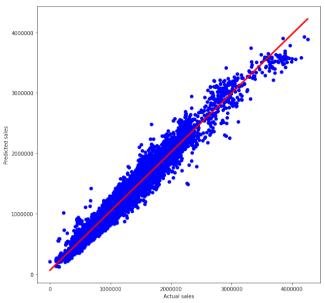


Figure 4.1: The visual description of the model shown to the users: a graph comparing the predicted sales and actual sales based on the original model. The red line depicts perfect predictions.

4.4.3 User Study Design

We test our method on a real dataset with real users, both from the company. We include a short tutorial about predictive modeling along with some questions to check users' understanding as a preliminary component of the study. This is because our users are a diverse set of individuals with a wide range of capabilities, including data scientists, human resource strategists, and senior members of the executive team. We also include participants from the University of Amsterdam to simulate users who could one day work in this environment. In total, we have 75 participants: 44 in the treatment group and 31 in the control group.

All users are first provided with a visual description of the model: a simple scatter plot comparing the predicted and actual sales (as shown in Figure 4.1). We also show a pie chart depicting the proportion of predictions that result in large errors to give users a sense of how frequently these mistakes occur. In our case, this is 4%. Since our users are diverse, we want to make our description of the model as accessible as possible while allowing them to form their own opinions about how well the model performs. Participants in the treatment group are shown MC-BRP explanations, while those in the control group are not given any explanation.

The study contains two components, objective and subjective, corresponding to RQ3.1 and RQ3.2, respectively. The objective component is meant to quantitatively evaluate whether or not users understand explanations generated by MC-BRP, while the subjective component assesses the effect of seeing the explanation on users' attitudes towards, and perceptions of, the model.

Type	Provide user with	User's task
Forward	(1) Input values(2) Explanation	Simulate output
Counterfactual	(1) Input values(2) Explanation(3) Output	Manipulate input to change output

Table 4.3: Summary of simulations performed in objective portion of the user study.

We base the objective component on *human-grounded metrics*, a framework proposed by Doshi-Velez and Kim [2018], where the tasks conducted by users are simplified versions of the original task. We modify the original sales prediction task into a binary classification one: we ask users to determine whether or not a prediction will result in a large error, as it seems unreasonable to expect humans to correctly predict retail sales values of order 10^6 .

Table 4.4: Summary of tasks performed in user study for the treatment and control groups. The subjective questions are asked twice.

Treatment	Control
Short modeling tutorial	Short modeling tutorial
Visual model description	Visual model description
Subjective questions	Subjective questions
Objective questions	Dummy questions
Subjective questions	Subjective questions

To answer **RQ3.1**, we ask users in the treatment group to perform two types of simulations, both suggested by Doshi-Velez and Kim [2018] and summarized in Table 4.3. The first is *forward simulation*, where we provide participants with the (i) input values, and (ii) explanation. We then ask them to simulate the output — whether or not this prediction will result in a large error. The second is *counterfactual simulation*, where we provide participants with the (i) input values, (ii) explanation, and (iii) output. We then ask them what they would have needed to change in the input in order to change the output. In other words, we want participants to determine how the input features can be changed (according to the trend) in order to produce a reasonable prediction as opposed to one that results in large error. These objective questions are designed to test whether or not a participant understands the explanations enough to predict or manipulate the model's output. We ask every participant in the treatment group to perform two forward simulations and one counterfactual simulation, and we show the same examples to all users.

For the control group, we found that we could not ask the objective questions in the same way we did for the treatment group. This is because the objective component involves simulating the model based on the explanations (see Table 4.3), which is not possible if the explanations are not provided. In fact, we initially left the objective questions in the control group study, but preliminary testing on some users from the company showed that this was confusing and unclear, similar to when we tried using LIME explanations. We were concerned this confusion would skew users' perceptions of the model and therefore convolute the results of **RQ3.2**. Instead, we show participants in the control group the (i) input values, and (ii) output – whether or not the example resulted in a large error. In this case, we ask them *if they have enough information* to determine why the example does (or does not) result in a large error. This serves as a dummy question to engage users with the task without confusing them. We cannot ask users in the control group to simulate the model since they do not see the explanations, but we want to mimic the conditions of the treatment group as closely as possible. Therefore, **RQ3.1**, is solely evaluated on users from the treatment group.

To answer **RQ3.2**, we contrast results from the treatment and control groups. We ask both groups of users the same four subjective questions twice, once towards the beginning of the study and once again at the end. We ask the questions at the beginning of the study to evaluate the distribution of preliminary attitudes towards the model, based solely on the visual description. We ask the questions at the end of the study to evaluate the effectiveness of MC-BRP explanations, by comparing the results from the treatment and control groups. The questions we devised are based on the user study by ter Hoeve et al. [2017]. Table 4.4 summarizes the experimental setup for the treatment and control groups. Again, the treatment and control groups are treated exactly the same with the exception of the objective questions – we only ask these to the treatment group since we cannot ask users to simulate the model without giving them the explanation.

4.5 EXPERIMENTAL RESULTS

In this section, we evaluate the explanations generated by MC-BRP in terms of (i) objective questions, and (ii) subjective questions.

4.5.1 Objective Questions

The results for users' objective comprehension of MC-BRP explanations are summarized in Table 4.5. We see that explanations generated by MC-BRP are both: (i) interpretable and (ii) actionable, with an average accuracy of 81.1%. This answers **RQ3.1**. When asked to perform forward simulations, the proportion of correct answers was 84.1% for both questions. This indicates that the majority of users were able to interpret the explanations in order to simulate the model's output (**RQ3.1**: interpretable). When asked to perform counterfactual simulations, the proportion of correct answers was slightly lower at 75.0%, but still indicates that the majority of users were able to determine how to manipulate the model's input in order to change the output (**RQ3.1**: actionable).

Human accuracy	
Forward simulation 1	84.1%
Forward simulation 2	84.1%
Counterfactual simulation	75.0%
Average	81.1%

Table 4.5: Results from the objective questions in the user study.

4.5.2 Subjective Questions

In order to understand the impact of MC-BRP explanations on users' attitudes towards the model, we ask them the following subjective questions:

- **SQ1:** I understand why the model makes large errors in predictions.
- SQ2: I would support using this model as a forecasting tool.
- SQ3: I trust this model.
- **SQ4**: In my opinion, this model produces mostly reasonable outputs.

To ensure our populations did not have different initial attitudes towards the model, we compared their answers on the subjective questions after only showing a visual description of the model. The visual description is a graph comparing the predicted sales to the actual sales, which allows users to see the distribution of errors made by the model (see Figure 4.1). We found no statistically significant difference (χ^2 test, $\alpha=0.05$) in initial attitudes towards the model, which allows us to postulate that any difference discovered between the two groups is a result of the treatment they were given (i.e., MC-BRP explanation vs. no explanation).

Figure 4.2 shows the distributions of answers to the four subjective questions in the treatment and control groups. The difference in distributions is significant for SQ1 ($\chi^2=18.2,\,\alpha=0.0001$): users in the treatment group agree with the statement more than users in the control group. However, we find no statistically significant difference between the two groups for the remaining questions (χ^2 test, $\alpha=0.05$). That is, MC-BRP explanations help users understand why the model makes large errors in predictions, but do not have an impact on users' trust or confidence in the model, or on their willingness to support its deployment.

4.6 DISCUSSION

Since our original motivation was to provide an explanation system that can be used by analysts at the company, we conducted a more in-depth analysis of the results to determine if there was a difference in attitudes between users depending on their background (e.g., Practitioners from the company or Researchers from the University of Amsterdam).

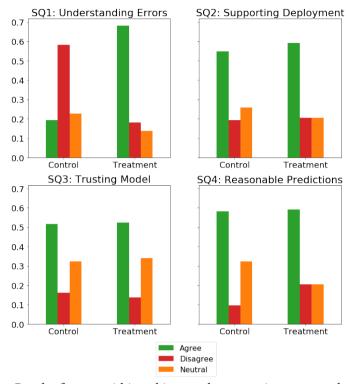


Figure 4.2: Results from a within-subject study comparing answers between the Treatment (MC-BRP explanation) and Control (no explanation) groups.

4.6.1 Comparing Attitudes Conditioned on Background

Table 4.6 shows the distribution of Practitioners and Researchers in the treatment and control groups. Since we have a slight imbalance in background between the treatment and control groups, we test whether or not our results still hold when conditioning on background and confirm that they do.

Again, we do not find statistically significant differences in initial attitudes towards the model (χ^2 test, $\alpha=0.05$). For Researchers, the distribution of answers between treatment and control groups is significantly different for SQ1 ($\chi^2=14.2, \alpha=0.001$), but does not differ for SQ2, SQ3, or SQ4 (χ^2 test, $\alpha=0.05$). The same holds for Practitioners: the distributions are significantly different only for SQ1 ($\chi^2=6.94, \alpha=0.05$). This is consistent with our results in Section 4.5. In both cases, users in the treatment group agree with SQ1 more than users in the control group, indicating that MC-BRP explanations help users understand why the model makes large errors in predictions, regardless of whether they are Practitioners or Researchers. Although the results are statistically significant for both groups, it should be noted that the results hold more strongly for Researchers compared to those for Practitioners, given the χ^2 values.

Table 4.6: Distribution of Practitioners and Researchers in the treatment and control groups.

Background	Practitioners	Researchers
Treatment	52%	48%
Control	58%	42%

4.6.2 Comparing Attitudes in the Treatment Group

Based on the users who saw the explanations, we compare the distributions of answers between Practitioners and Researchers in Figure 4.3 in order to understand the needs of different types of users. We find that there is a significant difference between Practitioners and Researchers for SQ2 ($\chi^2=7.94$, $\alpha=0.05$), indicating that more Resesearchers are in favor of using the model as a forecasting tool, and less are against it or have a neutral attitude, in comparison to the Practitioners. We also find a significant difference for SQ3 ($\chi^2=5.98$, $\alpha=0.05$): a larger proportion of Researchers trust the model, while the majority of Practitioners have neutral feelings. The results for SQ4 are significant as well ($\chi^2=6.86$, $\alpha=0.05$): although the majority of users in both groups believe the model produces reasonable predictions, a larger proportion of the Practitioners disagree with this statement in comparison to the Researchers.

We see no significant difference between groups for SQ1 (χ^2 test, $\alpha=0.05$), which makes sense given that we showed that MC-BRP explanations have a similar effect on both Practitioners and Researchers when comparing users in the treatment and control groups in Section 4.6.1.

Overall, these results suggest that our user study population is fairly heterogeneous, and that users from different backgrounds have different criteria for deploying or trusting a model, and varying levels of confidence regarding the accuracy of its outcomes.

4.6.3 User Study Limitations

Like any user study, ours has some limitations. It would have been preferable to distribute users more evenly in terms of the proportion of users in the treatment and control groups, as well as the proportion of Practitioners and Researchers in each of these groups. Unfortunately, this was not possible in our case because we recruited participants in two rounds: first for the treatment group, and then afterwards for the control group. One option could be to discard some Practitioners in the control group in order to have a better balance in terms of background, but we felt it was more important to have as many users as possible, and it would not be clear how to choose which users to discard. Fortunately, we found that our results still hold when conditioning on background as mentioned in Section 4.6.1. In future work, we plan to recruit for both groups at the same time to avoid issues like these.

We also acknowledge that not having a baseline method to compare to is a limitation of our study. In our case, the main issue is that there simply does not exist a method that is specifically for explaining errors in regression predictions, which would make asking questions about errors (i) unfair, and (ii) confusing, as mentioned

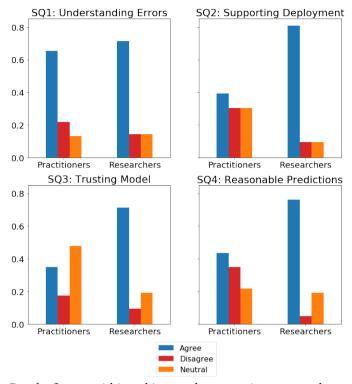


Figure 4.3: Results from a within-subject study comparing answers between participants who are Practitioners or Researchers (in the treatment group).

in Sections 4.4.2 and 4.4.3. However, now that MC-BRP exists, it can serve as a baseline for future work on erroneous predictions, which is another contribution of this paper.

4.7 CONCLUSION

In this chapter, we have proposed a method, MC-BRP, that provides users with contrastive explanations about predictions resulting in large errors based on: (i) the set of bounds for which reasonable predictions would be expected for each of the most important features. (ii) the trend between each of these features and the target.

Given a large error, MC-BRP generates a set of perturbed versions of the original instance that result in reasonable predictions. This is done by performing Monte Carlo simulations on each of the features deemed most important for the original prediction. For each of these features, we determine the bounds needed for a reasonable prediction based on the mean and standard deviation of this new set of reasonable predictions. We also determine the relationship between each feature and the target through the Pearson correlation, and present these to the user as the explanation.

We evaluate MC-BRP both objectively (**RQ3.1**) and subjectively (**RQ3.2**) by conducting a user study with 75 real users from the company and the University of

Amsterdam. We answer **RQ3.1** by conducting two types of simulations to quantify how (i) interpretable, and (ii) actionable our explanations are. Through forward simulations, we show that users are able to interpret MC-BRP explanations by simulating the model's output with an average accuracy of 84.5%. Through counterfactual simulations, we show that MC-BRP explanations are actionable with an accuracy of 76.2%.

We answer **RQ3.2** by conducting a between-subject experiment with subjective questions. The treatment group sees MC-BRP explanations, while the control group does not see any explanation. We find that explanations generated by MC-BRP help users understand why models make large errors in predictions (SQ1), but do not have a significant impact on support in deploying the model (SQ2), trust in the model (SQ3), or perceptions of the model's performance (SQ4). These results still hold when conditioning on users' background (Practitioners vs. Researchers). We also conduct an analysis on the treatment group to compare results between Practitioners and Researchers. We find significant differences for SQ2, SQ3 and SQ4, but do not find a significant difference in attitudes for SQ1.

This answers **RQ3**: we can create an explanation method based on a real-world use case by first identifying a use case where explanations are required and subsequently developing a method specific to this use case. We can evaluate in a human-centric manner by conducting a user study that includes both objective and subjective components.

For future work, we intend to explore allowing a predictive model to abstain from prediction when a particular instance has unusual feature values and determine the impact this has on users' trust, deployment support and perception of the model's performance. We also plan to compile a more comprehensive set of subjective questions by using multiple questions to evaluate users' impressions on the same topic.

So far, this thesis has focused on creating knowledge about responsible AI practices, specifically on developing new methods for explaining predictions from ML models. In the next and final part of the thesis, we will focus more on how to *translate* knowledge about responsible AI practices to the next generation of AI researchers.

REPRODUCIBILITY

To facilitate the reproducibility of the work in this chapter, our code is available at https://github.com/a-lucic/mcbrp.

Part III Pedagogy

Teaching Responsible AI through Reproducibility

In this part of the thesis, we investigate how to communicate responsible AI practices to the next generation of AI researchers. In this chapter, we address the following research question:

RQ4: How can we teach about responsible AI topics to a technical, research-oriented audience?

We answer **RQ4** by designing a course that is centered on a reproducibility project, where students work in teams to reimplement existing responsible AI algorithms from top AI conferences and reproduce the experiments reported in the papers.

5.1 INTRODUCTION

For several decades, the University of Amsterdam has offered a research-oriented Master of Science (MSc) program in AI. The main focus of the program is on the technical ML aspects of the major sub-fields of AI, such as computer vision, information retrieval, natural language processing, and reinforcement learning. One of the most recent additions to the MSc AI curriculum is a mandatory course on Fairness, Accountability, Confidentiality and Transparency in Artificial Intelligence (FACT-AI). This course was first taught during the 2019–2020 academic year and focuses on teaching FACT-AI topics through the lens of reproducibility. The main project involves students working in groups to re-implement existing FACT-AI algorithms from papers in top AI venues. There are approximately 150 students enrolled in the course each year.

The motivation for the course came from the MSc AI students themselves, who often play an important role in shaping the curriculum in order to meet the evolving requirements of researchers in both academia and industry. As the influence of AI on decision making is becoming increasingly prevalent in day-to-day life, there is a growing consensus that stakeholders who take part in the design or implementation of AI algorithms should reflect on the ethical ramifications of their work, including

This chapter was published at the AAAI Symposium on Educational Advances in Artificial Intelligence (EAAI-AAAI 2022) under the title "Reproducibility as a Mechanism for Teaching Fairness, Accountability, Confidentiality, and Transparency in Artificial Intelligence" [Lucic et al., 2022a].

developers and researchers [Campbell et al., 2021]. This is especially true in situations where data-driven AI systems affect some demographic sub-groups differently than others [Angwin et al., 2016; O'Neil, 2017]. As a result, our students have shown an increased interest in the ethical issues surrounding AI systems and requested that the university put together a new course focusing on responsible AI.

Since our MSc AI program is characterized by a strong emphasis on understanding, developing, and building AI algorithms, we believe that a new course on responsible AI in this program should also have a hands-on approach. The course is designed to address technical aspects of key areas in responsible AI: (i) fairness, (ii) accountability, (iii) confidentiality, and (iv) transparency, which we operationalize through a reproducibility project. We believe a strong emphasis on reproducibility is important from both an educational point of view and from the point of view of the AI community, since the (lack of) reproducible results has become a major point of critique in AI [Hutson, 2018]. Moreover, the starting point of almost any junior AI researcher (and most AI research projects in general) is re-implementing existing methods as baselines. The FACT-AI course is situated at a point in the program where students have learned the basics of ML and are ready to start experimenting with, and building on top of, state-of-the-art algorithms. Given that our MSc AI program is fairly research-oriented, it is important for students to experience the process of reproducing work done by others (and how difficult this is) at an early stage in their careers. We also believe reproducibility is a fundamental component of FACT-AI: the cornerstone of fair, accountable, confidential and transparent AI systems is having correct and reproducible results. Without reproducibility, it is unclear how to judge if a decision-making algorithm adheres to any of the FACT principles.

In the 2019–2020 academic year, we operationalized our learning ambitions regarding reproducibility by publishing a public repository with selected code implementations and corresponding reports from the group projects. In the 2020–2021 academic year, we took the projects one step further and encouraged students to submit to the ML Reproducibility Challenge, ¹ a competition that solicits reproducibility reports for papers published in conferences such as NeurIPS, ICML, ICLR, ACL, EMNLP, CVPR and ECCV. Although the MLRC broadly focuses on all papers submitted to these conferences, we focus exclusively on papers covering FACT-AI topics in our course. Submitting to the MLRC gives students a chance to experience the whole AI research pipeline, from running experiments, to writing rebuttals, to receiving the official notifications. Of the 23 papers that were accepted to the MLRC in 2021, 9 came from groups in the FACT-AI course at the University of Amsterdam.

In this chapter, we describe the *Fairness, Accountability, Confidentiality, and Transparency in Artificial Intelligence* course: a one month, full-time course based on examining ethical issues in AI using reproducibility as a pedagogical tool. Students work in groups to re-implement (and possibly extend) existing algorithms from top AI venues on FACT-AI topics. The course also includes lectures that cover the high-level principles of FACT-AI topics, as well as paper discussion sessions where students read and digest prominent FACT-AI papers. In this chapter, we outline the setup for the FACT-AI course and the experiences we had while running the course during the

¹https://paperswithcode.com/rc2020

2019-2020 and 2020-2021 academic years at the University of Amsterdam.

The remainder of this chapter is structured as follows. In Section 5.2 we discuss related work, specifically other courses about responsible AI. In Section 5.3, we detail ongoing reproducibility efforts in the AI community. In Section 5.4, we explain the learning objectives for our course, and explain how we realized those objectives in Section 5.5. We reflect on the feedback we received about the course in Section 5.6, as well as what worked (Section 5.7) and what did not (Section 5.8), before concluding in Section 5.9.

5.2 RELATED WORK

There have been multiple calls for introducing ethics in computer science courses in general, and in AI programs in particular [Angwin et al., 2016; Leonelli, 2016; O'Neil, 2017; of Sciences Engineering and Medicine, 2018; Singer, 2018; Skirpan et al., 2018; Grosz et al., 2019; Saltz et al., 2019; Danyluk et al., 2021]. Several surveys have investigated how existing responsible computing courses are organized [Peck, 2017; Fiesler et al., 2020; Garrett et al., 2020; Raji et al., 2021].

5.2.1 Characterizing Responsible AI Courses

There are two primary approaches to integrating such components into the curriculum: (i) stand-alone courses that focus on ethical issues such as FACT-AI topics, and (ii) a holistic curriculum where ethical issues are introduced and tackled in each course [Peck, 2017; Fiesler et al., 2020; Garrett et al., 2020]. In general, the latter is rare [Peck, 2017; Saltz et al., 2019; Fiesler et al., 2020], and can be difficult to organize due to a lack of qualified faculty or relevant expertise [Bates et al., 2020; Raji et al., 2021]. We opt for the first approach since our course is a new addition to an existing study program.

Fiesler et al. [2020] analyze 202 courses on "tech ethics". Their survey examines (i) the departments the courses are taught from, as well as the home departments of the course instructors, (ii) the topics covered in the courses, organized into 15 categories, and (iii) the learning outcomes in the courses. In our case, both the FACT-AI course and its instructors are from the Informatics Institute of the Faculty of Science at the University of Amsterdam. Our learning objectives (see Section 5.4) correspond to the following learning objectives from Fiesler et al. [2020]: "Critique", "Spot Issues", and "Create solutions". According to their content topic categorization, our course includes "AI & Algorithms" and "Research Ethics": the former since the course deals explicitly with AI algorithms under the FACT-AI umbrella, and the latter due to its focus on reproducibility. We note that "AI & Algorithms" is only the fifth-most popular topic according to the survey, after "Law & Policy", "Privacy & Surveillance", "Philosophy", and "Inequality, Justice & Human Rights" [see Table 2, Fiesler et al., 2020]). Although we believe these topics are important, we also wanted to avoid the feeling that the course was a "distraction from the real material" [Lewis and Stoyanovich, 2021], especially since (i) the majority of our students are coming from a technical background into a technical MSc program, and (ii) the FACT-AI course is mandatory for all students in the MSc AI program.

5.2.2 Similar Responsible AI Courses

The two courses that are the most similar to ours are those of Lewis and Stoyanovich [2021] and Yildiz et al. [2021].

Lewis and Stoyanovich [2021] describe a course for responsible data science. Similar to our course, they focus on the technical aspects of AI, involving lectures, readings, and a final project. However, their course differs from ours since the main project in their course is focused on examining the interpretability of an automated decision making system, while the main project in our course is focused on reproducibility.

Yildiz et al. [2021] describe a course based on reproducing experiments from AI papers, focusing on "low-barrier" reproducibility. Similar to our course, this course involves replicating a paper from scratch or reproducing the experiments using existing code, performing hyperparameter sweeps, and testing with new data or with variant algorithms. Another similarity is that they released a public repository of re-implemented algorithms,² which we also did for the 2019–2020 iteration of our course.³ However, their course differs from ours since theirs focuses on AI papers in general, while our course focuses exclusively on FACT-AI papers.

There are several courses that focus more on the philosophical or social science perspectives of AI ethics. Green [2021] describes an undergraduate AI ethics course that teaches computer science majors to analyse issues using different ethical approaches and how to incorporate these into an explicit ethical agent. Shen et al. [2021] introduce a toolkit in the form of "Value Cards" to inform students and practitioners about the social impacts of ML models through deliberation. Green and Crotts [2020] propose an approach to ethics education using "argument schemes" that summarize key ethical considerations for specialized domains such as healthcare or national defense. Furey and Martin [2018] introduce ethics concepts, primarily utilitarianism, into an existing AI course about autonomous vehicles by studying several variations of the Trolley Problem. Burton et al. [2018] teach ethics through science fiction stories complemented with philosophy papers, allowing students to reflect and debate difficult content without emotional or personal investment since the stories are not tied to "real" issues. Skirpan et al. [2018] describe an undergraduate course on human-centred computing which integrates ethical thinking throughout the design of computational systems. Unlike these courses, our course focuses more on the technical aspects of ethical AI. However, incorporating such non-technical perspectives is something we would like to do in future iterations of our course, perhaps through one of the mechanisms employed by some of these courses.

5.3 INTEGRATING REPRODUCIBILITY OF AI RESEARCH INTO THE FACT-AI COURSE

There have been several criticisms about the lack of reproducibility in AI research. Some have postulated that this is due to a combination of unpublished code and high sensitivity of training parameters [Hutson, 2018], while others believe the rapid

²https://reproducedpapers.org/

³https://github.com/uva-fact-ai-course/uva-fact-ai-course

rate of progress in ML research results in a lack of empirical rigor [Sculley et al., 2018]. Although typically well-intentioned, some papers may disguise speculation as explanation, obfuscate content behind math or language, and fail to attribute the correct sources of empirical gains [Lipton and Steinhardt, 2019].

Several efforts have been made to investigate and increase the reproducibility of AI research. In 2021, NeurIPS introduced a paper checklist including questions about reproducibility, along with a template for submitting source code as supplementary material [Beygelzimer et al., 2021]. The Association of Computing Machinery introduced a badging system that indicates how reproducible a paper is [ACM, 2019]. Papers with Code is an organization that provides links to official code repositories in arXiv papers [Stojnic, 2020]. It also hosts an annual ML Reproducibility Challenge: a community-wide effort to investigate the reproducibility of papers accepted at top AI conferences, which we incorporated into the 2020–2021 iteration of the FACT-AI course.

In an ideal scenario, reproducibility issues would be handled prior to publication [Sculley et al., 2018], but it can be difficult to catch such shortcomings in the review process due to the increasing number of papers submitted to AI conferences. Therefore, we believe it is of utmost importance that the next generation of AI researchers – including our own students – can (i) identify and (ii) avoid these pitfalls while conducting their own work. This, in combination with the fact that reproducibility is a fundamental component of responsible AI research, is why we opted to teach the FACT-AI course through the lens of reproducibility.

Our course is centered around a group project where students re-implement a recent FACT-AI algorithm from a top AI conference. This project has three components: (i) a reproducibility report, (ii) an associated code base, and (iii) a group presentation. In Section 5.5.3, we provide more details on the project and the outputs it resulted in.

5.4 LEARNING OBJECTIVES

In the FACT-AI course, we aim to make students aware of two types of responsibility: (i) towards society in terms of potential implications of their research, and (ii) towards the research community in terms of producing reproducible research. In this section, we outline the learning objectives for the FACT-AI course and explain how it fits within the context of the MSc AI program at the University of Amsterdam.

Table 5.1 shows the setup of the first year of the 2-year MSc AI program. Each semester at the University of Amsterdam is divided into three periods: two 8-week periods followed by one 4-week period. During an 8-week period, students follow two courses in parallel. During the 4-week period, they only follow a single course. The FACT-AI course takes place during the 4-week period at the end of the first semester, after students have taken Computer Vision 1, Machine Learning 1, Natural Language Processing 1 and Deep Learning 1. It is the only course students follow during this period, so we believe it is beneficial to have them focus on one main project – reproducing an existing FACT-AI paper. The learning objectives for the course are as follows:

• LO #1: Understanding FACT topics. Students can explain the major notions

Table 5.1: The first year of the MSc AI program at the University of Amsterdam, 2019–2020.

Course	Sem. 1	Sem. 2	EC
Computer Vision 1			6
Machine Learning 1			6
Natural Language Processing 1			6
Deep Learning 1			6
Fairness, Accountability, Confidentiality			6
and Transparency in AI			
Information Retrieval 1			6
Knowledge Representation and Reasoning			6
Elective 1			6
Elective 2			6
Elective 3			6

of fairness, accountability, confidentiality, and transparency that have been proposed in the literature, along with their strengths and weaknesses.

- LO #2: Understanding algorithmic harm. Students can explain, motivate, and distinguish the main types of algorithmic harm, both in general and in terms of concrete examples where AI is being applied.
- LO #3: Familiarity with FACT methods. Students are familiar with recent peer-reviewed algorithmic approaches to fairness, accountability, confidentiality, and transparency in the literature.
- LO #4: Reproducing FACT solutions. Students can assess the degree to
 which recent algorithmic solutions are effective, especially with respect to the
 claims made in the original papers, while understanding their limitations and
 shortcomings.

5.5 COURSE SETUP

The FACT-AI course is organized around (i) lectures, (ii) paper discussions, and (iii) a group project. It has had two iterations so far: the 2019–2020 iteration was taught in person, while the 2020–2021 iteration was taught online due to the COVID-19 pandemic. In this section, we detail how we realized the learning objectives from Section 5.4 and describe the challenges in adapting the course to an online format.

5.5.1 Lectures

To further the understanding of FACT-AI topics (LO1), we provide one general lecture for each of the 4 topics, along with a lecture specifically about reproducibility. Lectures are an opportunity for students to familiarize themselves with algorithmic harm (LO2). Students are encouraged to ask questions that lead to discussions about potential harm

done by applications of AI. This was more challenging in the second iteration of the course due to the online format, but we hope that facilitating such discussions will be more straightforward once we return to in-person classes.

In addition to the general lectures, we also include some guest lectures. These are used to either discuss specific types of algorithmic harm (LO2), examine specific FACT-AI algorithms in depth (LO3), or expand on the non-technical aspects of FACT-AI. Some examples of guest lectures include a lecture on AI accountability from a legal perspective by an instructor from the law department of the University of Amsterdam, and a lecture by two former FACT-AI students who explained how they turned their group project into an ICML 2021 workshop paper [Neely et al., 2021].

5.5.2 Paper Discussions

The goal of the paper discussion sessions is for students to learn about prominent FACT-AI methods (LO3), and learn to think critically about the claims made in the papers we discuss (LO4). Students first read a seminal FACT-AI paper on their own while trying to answer the following questions:

- What are the main claims of the paper?
- What are the research questions?
- Does the experimental setup make sense, given the research questions?
- What are the answers to the research questions? Are these supported by experimental evidence?

Once students have read the papers, they participate in smaller discussion sessions with their peers about their answers to the questions above. After each discussion session, all the groups are brought back together for a "dissection" session, where an instructor goes over the same seminal paper, giving an overview of the papers' strengths and weaknesses.

Each session was presented by a different instructor to show that there is no single way of examining a research paper, and that different researchers will bring different perspectives to their assessment of papers. The following papers were covered during the discussion sessions: [Hardt et al., 2016] on fairness; [Ribeiro et al., 2016a] on transparency; and [Abadi et al., 2016] on confidentiality.

5.5.3 Group Project

Reproduction of a FACT-AI paper

The purpose of the group project is to have students investigate the claims made by the authors of recent FACT-AI papers by diving into the details of the methods and their implementations. Using what they have learned from the paper discussion sessions, students work in groups to re-implement an existing FACT-AI algorithm from a top AI conference and re-run the experiments in the paper to determine the degree to which they are reproducible (LO4). If the code is already available, then they must extend the

method in some way. The project consisted of three deliverables: (i) a reproducibility report, (ii) an associated code base, and (iii) a group presentation.

In order to ensure the project is feasible, we select 10–15 papers in advance for groups to choose from. Our criteria for including papers is as follows:

- The paper is on a FACT-AI topic.
- At least one dataset in the paper is publicly available.
- Experiments can be run on a single GPU (which we provide access to).
- It is reasonable for a group of 3–4 MSc AI students to re-implement the paper within the timeframe of the course. In our case, students work on this project for one-month full-time.

To ease the load for our teaching assistants (TAs), we have several groups working on the same paper. We assign papers to TAs based on their interests by asking them to rank the set of candidate papers in advance. We also encourage them to suggest alternative papers provided they fit the criteria. The TAs read the papers before the course starts in order to ensure they have a sufficient, in-depth understanding of the work such that they can guide students through the project. This also serves as an extra feasibility check, to ensure that the papers are indeed a good fit for our course.

Each group writes a report about their efforts following the structure of a standard research paper (i.e., introduction, methodology, experiments, results, conclusion). They also include aspects specific to reproducibility such as explaining the difficulties of implementing certain components, as well as describing any communication they had with the original authors. In addition to the source code, students provide all results in a Jupyter notebook along with a file to install the required environment.

First Iteration: Contributing to an Open Source Repository

In the 2019–2020 iteration of the course, we created a public repository on GitHub, which contains a selection of the implementations done by our students: https://github.com/uva-fact-ai-course/uva-fact-ai-course. The TAs who assisted with the course decided which implementations to include and cleaned up the code so it all fit into one cohesive repository. This had multiple motivations. First, it taught students how to improve the reproducibility of their own work by releasing the code, while also giving them a sense of contributing to the open-source community. Second, a public repository can serve as a starting point for personal development in their future careers; companies often ask to see existing code or projects that prospective employees have worked on. Some students added the project to their CVs, while others wrote blog posts about their experiences, linking back to the repository.

⁴https://omarelb.github.io/self-explaining-neural-networks/

Second Iteration: The Machine Learning Reproducibility Challenge

In the 2020–2021 iteration of the course, we formally participated in the annual ML Reproducibility Challenge run by Papers with Code [Stojnic, 2020] in order to expose our students to the peer-review process. This gave students something to strive towards and offered perspectives beyond simply getting a grade for the project. Most importantly, it gave them the opportunity to experience the full research pipeline: (i) reading a technical paper, (ii) understanding a paper's strength and weaknesses, (iii) implementing (and perhaps also extending) the paper, (iv) writing up the findings, (v) submitting to a venue with a deadline, (vi) obtaining feedback, (vii) writing a rebuttal, and (viii) receiving the official notification. To encourage students to formally submit to the ML Reproducibility Challenge, we offered a 5% boost to their final grades if they submitted. Of the 32 groups in the FACT-AI course, 30 (94%) groups submitted their reproducibility reports to the ML Reproducibility Challenge, of which 9 groups had their papers accepted.

5.5.4 Taking the Course Online

The second iteration of the course was taught in January 2021, when the COVID-19 pandemic forced us to move classes and interactions online. Students made use of various channels to communicate: WhatsApp, Discord, and Slack. Canvas was the primary mode of communication between the instructors and the students, allowing students to ask questions and instructors to communicate various announcements.

Lectures were live, with frequent Q&A breaks to stimulate interactivity. Paper discussion sessions were organized in different online meeting subrooms where students discussed the papers together. This proved to be a challenge: while some subrooms had productive discussions, others struggled to get the conversation going.

The reproducibility project was more difficult to launch remotely. Since students had done online classes for their entire first semester, some struggled to find a group of fellow students to team up with, especially those coming from outside the MSc AI program. Overall, while we had various communication means, the lack of physical interaction due to COVID-19 slowed down our course organization.

5.6 FEEDBACK

In this section, we discuss the feedback we received about the course from the perspective of participating students (Section 5.6.1) and from the ML Reproducibility Challenge reviews (Section 5.6.2).

5.6.1 Feedback from Students

Both iterations of the course were evaluated using the standard evaluation procedure provided by the University of Amsterdam. However, only 16% of students filled out the evaluation form (23 out of 144) in the 2020–2021 iteration, potentially because the evaluation forms were administered online instead of in-person. According to the evaluation procedure at our university, this is not enough for a reliable quantitative

estimate of student satisfaction. Therefore, we focus on the 2019–2020 iteration when reporting student satisfaction statistics, since 53% of students filled out the form (79 out of 149) that year.

The vast majority of students were (very) satisfied with the course overall (67.8%). More specifically, students expressed satisfaction with the following dimensions:

Academic challenge: 75.2% were (very) satisfied

• Supervision: 76.9% were (very) satisfied

• Feedback: 81.3% were (very) satisfied

• Workload: 91.3% were (very) satisfied

• Level of the course: 79.7% were (very) satisfied

• Level of the report: 94.8% were (very) satisfied

• Level of the presentation: 96.6% were (very) satisfied

Table 5.2(a) shows some of the qualitative feedback we received from students. Based on this, we believe these high scores are mostly the result of the reproducibility project. Students enjoyed doing the project, especially due to the intensive supervision from our experienced TAs. The dimensions where we received the lowest scores were on the lectures and the final presentation, where only 30.6% and 30.2% were (very) satisfied with these aspects, respectively. This may be because we only provided four (high-level) lectures on each of the four topics, in order to give students as much time as possible to focus on the reproducibility project. However, it should be noted that the overall scores for these components were not poor, but average: 3.1/5 for lectures and 3.0/5 for the presentation.

5.6.2 Feedback from the ML Reproducibility Challenge

Of the 30 reproducibility reports submitted to the ML Reproducibility Challenge in the 2020–2021 iteration, 9 were accepted for publication in the ReScience Journal. In total, the ML Reproducibility Challenge accepted 23 reports, meaning that almost 40% of the reports accepted to the ML Reproducibility Challenge were from the University of Amsterdam.⁵

The reviews were mostly positive, with the general consensus being that most teams had gone beyond the general expectation of simply re-implementing the algorithm and re-running the experiments. Our TAs encouraged students to examine the generalizability of the work that was reproduced, either by trying new datasets or hyperparameters, or by performing ablation studies. Multiple reproducibility reports managed to question the results of the original papers with experimentally-supported claims. Importantly, some reviewers emphasized that these reproducibility studies were solid starting points for future research projects. For the reports that were rejected, the main critiques were that (i) only a fraction of the original work was

⁵https://openreview.net/group?id=ML_Reproducibility_Challenge/2020

reproduced, or (ii) no new insights were given. Some projects also had flaws in the experimental setup. See Table 5.2(b) for quotes from the ML Reproducibility Challenge reviews.

5.7 FACTORS CONTRIBUTING TO A SUCCESSFUL COURSE

Understanding and re-implementing the work of other researchers is not a trivial task, especially for first-year MSc students. There were several aspects of the setup that we believe were beneficial for the students, which we organize along three dimensions: (i) general, (ii) concerning FACT-AI, and (iii) concerning reproducibility. We believe each of these factors is important for a successful implementation of this course, or other similar courses.

5.7.1 General

Timing of the course

It is important that students have prior knowledge of ML theory as well as some programming experience before completing a project-based course in groups. At the University of Amsterdam, the FACT-AI course takes place after students have completed 4 ML-focused courses (see Table 5.1). We believe it is extremely important that students have access to adequate preparation, especially in terms of programming experience, before setting off to reproduce experiments from prominent AI papers. Without this prior knowledge, we believe such a project would not be feasible in the allotted time frame.

Regular contact with experienced TAs

The TAs are there to help with two main components: (i) understanding the paper, and (ii) debugging the implementation process. In practice, we found that it is extremely important for the TAs to have excellent programming experience since this is the main aspect students need help with. We also had a dedicated Slack workspace for the TAs and course instructors to keep in touch regularly.

Since our course is only four weeks long, we found it was important for students to have regular contact with their TAs to ensure no one got stuck in the process. For the first (in-person) iteration of the course, groups had one-hour tutorials with their TAs twice per week, where all groups that were working on the same paper (and therefore had the same TA) were in the same tutorial. Since they were all working on the same paper, there were many overlapping questions, and students found it beneficial to be able to share their experiences with one another. For the second (online) iteration of the course, we thought it would be challenging to ensure each group got the attention they needed if everyone was in one large online tutorial, so the TAs met with each group separately for 30 minutes, twice per week.

Table 5.2: Feedback about the course.

(a) Feedback from students

- "Reproducing an article was hard and intensive but a really good experience."
- "Replicating another study, seeing how (poorly) other research is performed was really eye-opening."
- "Reproducing a paper: I believe this is a good thing to do and is an important part of academia."
- "Gave good insights into the trustworthiness of research papers, which is apparently not great."
- "I appreciate the critical view I have developed on papers as a result of this course. Normally I would easily accept the content of a paper, but I will be more critical from now on, as many papers are not reproducible."
- "I think it's really good that we get some practical insights into reproducing results from other papers, not all papers are as good as they seem to be."
- "I really appreciated that this was the first course where students are judging stateof-the-art AI-models. In other words, students were able to experience the scientific workfield of AI."

(b) Feedback from the ML Reproducibility Challenge

- "The report reveals a lot of dark spots of the original paper."
- "Good reviews, strong reproducibility report, provides code reimplementation from scratch which is a strong contribution."
- "The discussion section is a great reference point for future work."
- "The additional experimentation is rather impressive and the report reflects an intuitive understanding of concepts such as coverage, correctness, and counterfactual explanations."
- "The report provides good insights on how the experiments in the original paper actually work, while also generating new hypothesis to be tested for future research, which is a positive outcome."
- "My main concern is that it remains unclear why some of the results are so far off from the original paper? I would have expected the authors to dig deeper on that."
- "It doesn't go above and beyond the reproduction and does not offer novel insights into the workings of the original paper."
- "The submission failed at reproducing the original results. It is unclear whether this is due to a difference in the experimental setup or due to implementation errors."

Early feedback on the reports

Approximately halfway through the course, we asked students to submit a draft report to their TAs in order to get feedback. We found this significantly increased the quality of the final reports.

5.7.2 Concerning FACT-AI

Emphasizing the technical perspective of FACT-AI

Given that the FACT-AI course is situated in the context of a technical, research-oriented MSc, having students re-implement research papers from top AI conferences was an effective way to teach FACT-AI topics for our students. Teaching FACT-AI from a primarily technical perspective aligns well with what students expect from the MSc AI program at the University of Amsterdam. Although we believe a technical focus makes sense for our MSc program, we also believe it is important to incorporate non-technical perspectives into the course – see Section 5.8.2.

Creating resources for the FACT-AI community

We believe a significant motivating factor for students was creating concrete output that extended beyond simply completing a project for a course: creating resources for the FACT community. In the 2019–2020 iteration, this was done by creating a public repository with the best FACT-AI algorithm implementations, as selected by the TAs. In the 2020–2021 iteration, this was done by publicly submitting their reproducibility reports about FACT-AI algorithms to the ML Reproducibility Challenge, where the accepted reports were published in the ReScience Journal. In the future, we plan to continue aligning our course with the ML Reproducibility Challenge since we found the process extremely beneficial for our students.

5.7.3 Concerning Reproducibility

Including extension as part of reproducibility

If source code was already available for the paper – which is fortunately becoming increasingly common for AI research papers – we asked students to think about how to extend the paper since the implementation was already available. This resulted in some creative and interesting ideas in the reports, and we believe this is why our students performed well at the ML Reproducibility Challenge (see Section 5.6.2)

Simple grading setup.

For a 4-week, project-based course, we found it was beneficial for students to focus on one main deliverable consisting of three components: (i) the reproducibility report, (ii) the associated code base, and (iii) the group presentation. The report that students submitted for the course was the same one they submitted to the ML Reproducibility Challenge. This way, participating in the ML Reproducibility Challenge was not an extra task but rather an integral part of the course.

5.8 AREAS OF IMPROVEMENT

Although we believe both iterations of the course went well, there are several aspects of the setup that we believe could use some improvement and other instructors should consider if they plan to implement a similar course.

5.8.1 General

Given that this is the first time most students are formally submitting a paper, it is not surprising that there were some logistical issues. Some groups made minor mistakes such as forgetting to submit their work double-blind or slightly missing the submission deadline. We also had some groups who wrote the introduction sections of their papers as an introduction to the FACT-AI course, rather than an introduction to the topic they were working on. In future iterations, we will explicitly state the standard procedures of writing and submitting a paper and provide some examples.

5.8.2 Concerning FACT-AI

Although focusing primarily on the technical aspects of FACT-AI is an effective way to engage our technical students in socially-relevant AI problems, we also believe that they would benefit from additional non-technical perspectives on FACT-AI topics. In the future, we plan to include perspective from outside of computer science through (i) additional guest lectures, (ii) workshop sessions [Skirpan et al., 2018; Shen et al., 2021], and (iii) broader impact statements [Campbell et al., 2021] in the reproducibility reports.

5.8.3 Concerning Reproducibility

In future iterations, we believe it would be useful to show students more examples of what a high-quality reproducibility paper looks like and explain in-depth why it is high-quality. These could be papers that were previously accepted to the ML Reproducibility Challenge, or papers from other reproducibility efforts outlined in Section 5.3. We want the students to understand what makes a paper a good (reproducibility) paper, that is, it has a set of (reproducibility) claims, it argues for these claims, and shows evidence to support these claims.

5.9 CONCLUSION

In this chapter, we share our setup for the FACT-AI course at the University of Amsterdam, which teaches FACT-AI topics through reproducibility. The course set out to give students (i) an understanding of FACT-AI topics, (ii) an understanding of algorithmic harm, (iii) familiarity with recent FACT-AI methods, and (iv) an opportunity to reproduce FACT-AI solutions, through a combination of lectures, paper discussion sessions and a reproducibility project.

Through their projects, our students engaged with the open-source community by creating a public code repository (in the 2019–2020 iteration), as well as with the research community via successful submissions to the ML Reproducibility Challenge challenge (in the 2020–2021 iteration). We also detail how the 2020–2021 iteration brought about its own unique set of challenges due to the COVID-19 pandemic.

In this course, we illustrate that reproducibility is not only paramount to good science in general, but is also a fundamental component of FACT-AI. We received very positive feedback on teaching FACT-AI topics through reproducibility. We believe this was an excellent fit for our students, which not only helped motivate them for the duration of the course, but also helped them develop skills that will be essential in their future research careers, whether in the private or public sector. To generalize this course setup to other scientific domains, we suggest identifying where the lack of reproducibility in this domain area is coming from and centre the project around evaluating this component.

With this final chapter, we answer **RQ4**: we can use reproducibility as a mechanism for teaching responsible AI concepts to a technical, research-oriented audience. Structuring the course around a reproducibility project gives students the opportunity to learn about responsible AI concepts, such as explainability, in a hands-on manner.

6 Conclusions

In this chapter you look back at the questions you asked in Section 1.1 and formulate the answers based on the research chapters. You than continue with a section on future research directions that follow from the work in the thesis.

6.1 WHAT HAVE WE DONE?

rq:wasabi! rq:wasabi!

rq:stillwasabi! rq:stillwasabi!

6.2 WHAT IS NEXT?

Bibliography

- Martin Abadi, Andy Chu, Ian Goodfellow, H. Brendan McMahan, Ilya Mironov, Kunal Talwar, and Li Zhang. Deep Learning with Differential Privacy. In ACM SIGSAC Conference on Computer and Communications Security, 2016. (Cited on page 71.)
- ACM. Artifact Review and Badging. https://www.acm.org/publications/policies/artifact-review-badging, 2019. (Cited on page 69.)
- Ahmed M. Alaa, Michael Weisz, and Mihaela van der Schaar. Deep Counterfactual Networks with Propensity-Dropout. In ICML Workshop on Principled Approaches to Deep Learning, June 2017. (Cited on page 9.)
- Julia Angwin, Jeff Larson, Surya Mattu, and Lauren Kirchner. Machine Bias. ProPublica. https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing, 2016. (Cited on pages 66 and 67.)
- Federico Baldassarre and Hossein Azizpour. Explainability Techniques for Graph Convolutional Networks. In *ICML Workshop on Learning and Reasoning with Graph-Structured Representations*, 2019. (Cited on pages 28 and 29.)
- Randall Balestriero. Neural Decision Trees. arXiv, 2017. (Cited on page 10.)
- Solon Barocas, Andrew D. Selbst, and Manish Raghavan. The Hidden Assumptions Behind Counterfactual Explanations and Principal Reasons. In *ACM Conference on Fairness, Accountability, and Transparency*, 2020. (Cited on page 40.)
- Jo Bates, David Cameron, Alessandro Checco, Paul Clough, Frank Hopfgartner, Suvodeep Mazumdar, Laura Sbaffi, Peter Stordy, and Antonio de la Vega de León. Integrating FATE/Critical Data Studies into Data Science Curricula: Where Are We Going and How Do We Get There? In ACM Conference on Fairness, Accountability, and Transparency, 2020. (Cited on page 67.)
- Peter W. Battaglia, Jessica B. Hamrick, Victor Bapst, Alvaro Sanchez-Gonzalez, Vinicius Zambaldi, Mateusz Malinowski, Andrea Tacchetti, David Raposo, Adam Santoro, Ryan Faulkner, Caglar Gulcehre, Francis Song, Andrew Ballard, Justin Gilmer, George Dahl, Ashish Vaswani, Kelsey Allen, Charles Nash, Victoria Langston, Chris Dyer, Nicolas Heess, Daan Wierstra, Pushmeet Kohli, Matt Botvinick, Oriol Vinyals, Yujia Li, and Razvan Pascanu. Relational Inductive Biases, Deep Learning, and Graph Networks. *arXiv*, 2018. (Cited on page 30.)
- Alina Beygelzimer, Yann Dauphin, Percy Liang, and Jennifer Wortman-Vaughan. Introducing the NeurIPS 2021 Paper Checklist. https://neuripsconf.medium.com/introducing-the-neurips-2021-paper-checklist, 2021. (Cited on page 69.)
- Battista Biggio and Fabio Roli. Wild Patterns: Ten Years After the Rise of Adversarial Machine Learning. Pattern Recognition, 2018. (Cited on page 9.)
- Francesco Bodria, Fosca Giannotti, Riccardo Guidotti, Francesca Naretto, Dino Pedreschi, and Salvatore Rinzivillo. Benchmarking and Survey of Explanation Methods for Black Box Models. *arXiv*, 2021. (Cited on page 28.)
- Markus M. Breunig, Hans-Peter Kriegel, Raymond T. Ng, and Jörg Sander. LOF: Identifying Density-Based Local Outliers. ACM SIGDMOD Record, 2000. (Cited on page 19.)
- Tom B. Brown, Dandelion Mané, Aurko Roy, Martín Abadi, and Justin Gilmer. Adversarial Patch. In Advances in Neural Information Processing Systems, 2018. (Cited on page 9.)
- Emanuelle Burton, Judy Goldsmith, and Nicholas Mattei. How to Teach Computer Ethics through Science Fiction. *Communications of the ACM*, 2018. (Cited on page 68.)
- Rosie Campbell, Madhulika Srikumar, and Hudson Hongo. Managing the Risks of AI Research: Six Recommendations for Responsible Publication. Partnership on AI, https://www.partnershiponai.org/wp-content/uploads/2021/05/PAI-Managing-the-Risks-of-AI-Research-Responsible-Publication.pdf, 2021. (Cited on pages 66 and 78.)
- Ines Chami, Sami Abu-El-Haija, Bryan Perozzi, Christopher Ré, and Kevin Murphy. Machine Learning on Graphs: A Model and Comprehensive Taxonomy. arXiv, 2021. (Cited on page 30.)
- Andrea Danyluk, Paul Leidig, Andrew McGettrick, Lillian Cassel, Maureen Doyle, Christian Servin, Karl Schmitt, and Andreas Stefik. Computing Competencies for Undergraduate Data Science Programs: An ACM Task Force Final Report. In ACM Technical Symposium on Computer Science Education, 2021. (Cited on page 67.)
- Kim de Bie, Ana Lucic, and Hinda Haned. To Trust or Not to Trust a Regressor: Estimating and Explaining Trustworthiness of Regression Predictions. In ICML Workshop on Human in the Loop Learning, 2021.
- Andreea Deac, Yu-Hsiang Huang, Petar Veličković, Pietro Liò, and Jian Tang. Drug-Drug Adverse Effect Prediction with Graph Co-Attention. In ICML Workshop on Computational Biology, 2019. (Cited on page 28.)

- A.K. Debnath, R.L. Lopez de Compadre, G. Debnath, A.J. Shusterman, and C. Hansch. Structure-Activity Relationship of Mutagenic Aromatic and Heteroaromatic Nitro Compounds. Correlation with Molecular Orbital Energies and Hydrophobicity. *Journal of Medicinal Chemistry*, 1991. (Cited on page 35.)
- Amit Dhurandhar, Pin-Yu Chen, Ronny Luss, Chun-Chen Tu, Paishun Ting, Karthikeyan Shanmugam, and Payel Das. Explanations based on the Missing: Towards Contrastive Explanations with Pertinent Negatives. In *Advances in Neural Information Processing Systems*, 2018. (Cited on pages 9 and 49.)
- B. J. Dietvorst, J. P. Simmons, and C. Massey. Algorithm Aversion: People Erroneously Avoid Algorithms After Seeing Them Err. *Journal of Experimental Psychology*, 2015. (Cited on pages 45 and 46.)
- Finale Doshi-Velez and Been Kim. Considerations for Evaluation and Generalization in Interpretable Machine Learning. Springer International Publishing, 2018. (Cited on pages 7, 35, 45, and 55.)
- Alexandre Duval and Fragkiskos D. Malliaros. GraphSVX: Shapley Value Explanations for Graph Neural Networks. In European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases, 2021. (Cited on pages 28 and 29.)
- Upol Ehsan, Philipp Wintersbergerand Q. Vera Liao, Martina Mara, Marc Streit, Sandra Wachter, Andreas Riener, and Mark O. Riedl. Operationalizing Human-Centered Perspectives in Explainable AI. Workshop at CHI 2021., 2021. (Cited on page 40.)
- EU. Regulation (EU) 2016/679 of the European Parliament (GDPR). Official Journal of the European Union, 2016. (Cited on pages 7, 27, and 46.)
- Lukas Faber, Amin K Moghaddam, and Roger Wattenhofer. Contrastive Graph Neural Network Explanation. In *ICML Workshop on Graph Representation Learning and Beyond*, 2020. (Cited on page 29.)
- FICO. Explainable Machine Learning Challenge, 2017. URL https://community.fico.com/s/explainable-machine-learning-challenge. (Cited on pages 14 and 46.)
- Casey Fiesler, Natalie Garrett, and Nathan Beard. What Do We Teach When We Teach Tech Ethics? A Syllabi Analysis. In ACM Technical Symposium on Computer Science Education, 2020. (Cited on page 67.)
- Timo Freiesleben. The Intriguing Relation Between Counterfactual Explanations and Adversarial Examples. Minds and Machines, 2021. (Cited on page 30.)
- Heidi Furey and F. Martin. Introducing Ethical Thinking About Autonomous Vehicles Into an AI Course. In AAAI Symposium on Educational Advances in Artificial Intelligence, 2018. (Cited on page 68.)
- Natalie Garrett, Nathan Beard, and Casey Fiesler. More Than "If Time Allows": The Role of Ethics in AI Education. In AAAI/ACM Conference on AI, Ethics, and Society, 2020. (Cited on page 67.)
- R. Goebel, A. Chander, K. Holzinger, F. Lecue, Z. Akata, S. Stumpf, P. Kieseberg, and A. Holzinger. Explainable AI: The New 42? In *International IFIP Cross Domain Conference for Machine Learning and Knowledge Extraction*, 2018. (Cited on page 27.)
- Ian J. Goodfellow, Jonathon Shlens, and Christian Szegedy. Explaining and Harnessing Adversarial Examples. In International Conference on Learning Representations, 2015. (Cited on page 9.)
- Rory Mc Grath, Luca Costabello, Chan Le Van, Paul Sweeney, Farbod Kamiab, Zhao Shen, and Freddy Lecue. Interpretable Credit Application Predictions With Counterfactual Explanations. In *NeurIPS Workshop on Challenges and Opportunities for AI in Financial Services*, 2018. (Cited on page 9.)
- Nancy Green. An Al Ethics Course Highlighting Explicit Ethical Agents. In AAAI/ACM Conference on AI, Ethics, and Society, 2021. (Cited on page 68.)
- Nancy L Green and L Joshua Crotts. Argument Schemes for AI Ethics Education. In COMMA Workshop on Computational Models of Natural Argument, 2020. (Cited on page 68.)
- Barbara J. Grosz, David Gray Grant, Kate Vredenburgh, Jeff Behrends, Lily Hu, Alison Simmons, and Jim Waldo. Embedded EthiCS: Integrating Ethics across CS Education. *Communications of the ACM*, 2019. (Cited on page 67.)
- Riccardo Guidotti, Anna Monreale, Salvatore Ruggieri, Dino Pedreschi, Franco Turini, and Fosca Giannotti. Local Rule-Based Explanations of Black Box Decision Systems. *arXiv*, 2018a. (Cited on page 30.)
- Riccardo Guidotti, Anna Monreale, Franco Turini, Dino Pedreschi, and Fosca Giannotti. A Survey of Methods for Explaining Black Box Models. *ACM Computing Surveys*, 2018b. (Cited on pages 28, 48, and 52.)
- Zhichun Guo, Chuxu Zhang, Wenhao Yu, John Herr, Olaf Wiest, Meng Jiang, and Nitesh V. Chawla. Few-Shot Graph Learning for Molecular Property Prediction. In *The Web Conference*, 2021. (Cited on page 28.)
- Moritz Hardt, Eric Price, and Nathan Srebro. Equality of Opportunity in Supervised Learning. In *Advances in Neural Information Processing Systems*, 2016. (Cited on page 71.)
- Lisa Anne Hendricks, Ronghang Hu, Trevor Darrell, and Zeynep Akata. Grounding Visual Explanations. In *European Conference on Computer Vision*. 2018. (Cited on page 49.)
- Dennis J. Hilton. Conversational Processes and Causal Explanation. *Psychological Bulletin*, 1990. (Cited on pages 50 and 52.)
- Dennis J. Hilton. Social Attribution and Explanation. The Oxford Handbook of Causal Reasoning, 2017. (Cited

- on page 53.)
- Dennis J. Hilton and Ben R. Slugoski. Knowledge-based Causal Attribution: The Abnormal Conditions Focus Model. *Psychological Review*, 1986. (Cited on pages 46 and 50.)
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the Knowledge in a Neural Network. In NeurIPS Workshop on Deep Learning, March 2014. (Cited on page 10.)
- Qiang Huang, Makoto Yamada, Yuan Tian, Dinesh Singh, Dawei Yin, and Yi Chang. GraphLIME: Local Interpretable Model Explanations for Graph Neural Networks. *arXiv*, 2020. (Cited on page 29.)
- Matthew Hutson. Artificial Intelligence Faces Reproducibility Crisis. Science, 2018. (Cited on pages 66 and 68.)
- IJCAI. Explainable AI Workshop, 2017. (Cited on page 49.)
- Shalmali Joshi, Oluwasanmi Koyejo, Warut Vijitbenjaronk, Been Kim, and Joydeep Ghosh. Towards Realistic Individual Recourse and Actionable Explanations in Black-Box Decision Making Systems. In *International Conference on Artificial Intelligence and Statistics*, 2020. (Cited on page 9.)
- Kentaro Kanamori, Takuya Takagi, Ken Kobayashi, and Hiroki Arimura. DACE: Distribution-Aware Counterfactual Explanation by Mixed-Integer Linear Optimization. In *International Joint Conference on Artificial Intelligence*, 2020. (Cited on pages 9, 14, 15, 19, and 30.)
- Bo Kang, Jefrey Lijffijt, and Tijl De Bie. ExplaiNE: An Approach for Explaining Network Embedding-based Link Predictions. *arXiv*, 2019. (Cited on page 29.)
- Amir-Hossein Karimi, Gilles Barthe, Borja Balle, and Isabel Valera. Model-Agnostic Counterfactual Explanations for Consequential Decisions. In *International Conference on Artificial Intelligence and Statistics*, 2020a. (Cited on pages 9, 10, and 28.)
- Amir-Hossein Karimi, Gilles Barthe, Bernhard Schölkopf, and Isabel Valera. A Survey of Algorithmic Recourse: Definitions, Formulations, Solutions, and Prospects. *arXiv*, 2020b. (Cited on pages 28 and 30.)
- Amir-Hossein Karimi, Julius von Kügelgen, Bernhard Schölkopf, and Isabel Valera. Algorithmic Recourse Under Imperfect Causal Knowledge: A Probabilistic Approach. In *Advances in Neural Information Processing Systems*, 2020c. (Cited on page 9.)
- Amir-Hossein Karimi, Bernhard Scholkopf, and Isabel Valera. Algorithmic Recourse: From Counterfactual Explanations to Interventions. In ACM Conference on Fairness, Accountability, and Transparency, 2021. (Cited on page 9.)
- Surya Karunagaran, Ana Lucic, and Christine Custis. XAI Toolsheets: An Evaluation Framework for XAI Tools. In *Under review at AAAI/ACM Conference on AI, Ethics, and Society*, 2022.
- Atoosa Kasirzadeh and Andrew Smart. The Use and Misuse of Counterfactuals in Ethical Machine Learning. In ACM Conference on Fairness, Accountability, and Transparency, 2021. (Cited on page 40.)
- Amir E. Khandani, Adlar J. Kim, and Andrew W. Lo. Consumer Credit-risk Models via Machine Learning Algorithms. *Journal of Banking and Finance*, 2010. (Cited on page 46.)
- Diederik P. Kingma and Jimmy Ba. Adam: A Method for Stochastic Optimization. In *International Conference on Learning Representations*, 2015. (Cited on page 14.)
- Thomas N. Kipf and Max Welling. Semi-Supervised Classification with Graph Convolutional Networks. In *International Conference on Learning Representations*, 2017. (Cited on page 32.)
- Thibault Laugel, Marie-Jeanne Lesot, Christophe Marsala, Xavier Renard, and Marcin Detyniecki. Inverse Classification for Comparison-based Interpretability in Machine Learning. In *International Conference on Information Processing and Management of Uncertainty in Knowledge-Based Systems*, 2018. (Cited on pages 9, 10, and 30.)
- Sabina Leonelli. Locating Ethics in Data Science: Responsibility and Accountability in Global and Distributed Knowledge Production Systems. Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences, 2016. (Cited on page 67.)
- Armanda Lewis and Julia Stoyanovich. Teaching Responsible Data Science: Charting New Pedagogical Territory. *International Journal of Artificial Intelligence in Education*, 2021. (Cited on pages 67 and 68.)
- Wanyu Lin, Hao Lan, and Baochun Li. Generative Causal Explanations for Graph Neural Networks. In *International Conference on Machine Learning*, 2021. (Cited on pages 28 and 29.)
- Peter Lipton. Contrastive Explanation. Royal Institute of Philosophy Supplement, 1990. (Cited on pages 50 and 52.)
- Zachary C Lipton and Jacob Steinhardt. Research for Practice: Troubling Trends in Machine-Learning Scholarship. *Communications of the ACM*, 2019. (Cited on page 69.)
- Ana Lucic, Hinda Haned, and Maarten de Rijke. Explaining Predictions from Tree-based Boosting Ensembles. In SIGIR Workshop on Fairness, Accountability, Confidentiality, and Transparency in Information Retrieval, 2019.
- Ana Lucic, Hinda Haned, and Maarten de Rijke. Why Does My Model Fail? Contrastive Local Explanations

- for Retail Forecasting. In ACM Conference on Fairness, Accountability, and Transparency, 2020. (Cited on pages 30 and 45.)
- Ana Lucic, Madhulika Srikumar, Umang Bhatt, Alice Xiang, Ankur Taly, Q. Vera Liao, and Maarten de Rijke. A Multistakeholder Approach Towards Evaluating AI Transparency Mechanisms. In *CHI Workshop on Operationalizing Human-Centered Perspectives in Explainable AI*, 2021.
- Ana Lucic, Maurits Bleeker, Sami Jullien, Samarth Bhargav, and Maarten de Rijke. Reproducibility as a Mechanism for Teaching Fairness, Accountability, Confidentiality, and Transparency in Artificial Intelligence. In AAAI Symposium on Educational Advances in Artificial Intelligence, 2022a. (Cited on page 65.)
- Ana Lucic, Harrie Oosterhuis, Hinda Haned, and Maarten de Rijke. FOCUS: Flexible Optimizable Counterfactual Explanations for Tree Ensembles. In *AAAI Conference on Artificial Intelligence*, 2022b. (Cited on pages 7, 30, and 31.)
- Ana Lucic, Maartje ter Hoeve, Gabriele Tolomei, Maarten de Rijke, and Fabrizio Silvestri. CF-GNNExplainer: Counterfactual Explanations for Graph Neural Networks. In *International Conference on Artificial Intelligence and Statistics*, 2022c. (Cited on page 27.)
- Scott M Lundberg and Su-In Lee. A Unified Approach to Interpreting Model Predictions. In *Advances in Neural Information Processing Systems*, 2017. (Cited on pages 29, 34, 49, and 50.)
- Scott M. Lundberg, Gabriel Erion, Hugh Chen, Alex DeGrave, Jordan M. Prutkin, Bala Nair, Ronit Katz, Jonathan Himmelfarb, Nisha Bansal, and Su-In Lee. From Local Explanations to Global Understanding with Explainable AI for Trees. *Nature Machine Intelligence*, 2020. (Cited on pages 48 and 49.)
- Dongsheng Luo, Wei Cheng, Dongkuan Xu, Wenchao Yu, Bo Zong, Haifeng Chen, and Xiang Zhang. Parameterized Explainer for Graph Neural Network. In *Advances in Neural Information Processing Systems*, 2020. (Cited on pages 28, 29, 35, and 37.)
- Prashan Madumal, Tim Miller, Liz Sonenberg, and Frank Vetere. Explainable Reinforcement Learning Through a Causal Lens. In *AAAI Conference on Artificial Intelligence*, 2019. (Cited on page 9.)
- Tim Miller. Explanation in Artificial Intelligence: Insights from the Social Sciences. *Artificial Intelligence*, 2019. (Cited on pages 1, 27, and 52.)
- Tim Miller, Piers Howe, and Liz Sonenberg. Explainable AI: Beware of Inmates Running the Asylum Or: How I Learnt to Stop Worrying and Love the Social and Behavioural Sciences. In *IJCAI Workshop on Explainable AI*, 2017. (Cited on page 49.)
- Christoph Molnar. Interpretable Machine Learning: A Guide for Making Black Box Models Explainable. Lulu.com, 2019. (Cited on page 36.)
- Ramaravind K. Mothilal, Amit Sharma, and Chenhao Tan. Explaining machine learning classifiers through diverse counterfactual explanations. In *ACM Conference on Fairness, Accountability, and Transparency*, 2020. (Cited on page 9.)
- Michael Neely, Stefan F. Schouten, Maurits J. R. Bleeker, and Ana Lucic. Order in the Court: Explainable AI Methods Prone to Disagreement. In *ICML Workshop on Theoretic Foundation, Criticism, and Application Trend of Explainable AI*, 2021. (Cited on page 71.)
- Cuong Q. Nguyen, Constantine Kreatsoulas, and Kim M. Branson. Meta-Learning GNN Initializations for Low-Resource Molecular Property Prediction. In *ICML Workshop on Graph Representation Learning and Beyond*, 2020. (Cited on page 28.)
- National Academies of Sciences Engineering and Medicine. Data Science for Undergraduates: Opportunities and Options. National Academies Press, 2018. (Cited on page 67.)
- Dan Ofer. COMPAS Dataset. https://www.kaggle.com/danofer/compass, 2017. (Cited on pages 14 and 15.)
- Cathy O'Neil. The Ivory Tower Can't Keep Ignoring Tech. The New York Times: https://www.nytimes.com/2017/11/14/opinion/academia-tech-algorithms.html, 2017. (Cited on pages 66 and 67.)
- Evan M. Peck. The Ethical Engine: Integrating Ethical Design into Intro Computer Science. Medium, https://medium.com/bucknell-hci/the-ethical-engine-integrating-ethical-design-into-intro-to-computer-science-4f9874e756af, 2017. (Cited on page 67.)
- Marco A.F. Pimentel, David A. Clifton, Lei Clifton, and Lionel Tarassenko. A Review of Novelty Detection. Signal Processing, 2014. (Cited on page 49.)
- Phillip E. Pope, Soheil Kolouri, Mohammad Rostami, Charles E. Martin, and Heiko Hoffmann. Explainability Methods for Graph Convolutional Neural Networks. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019. (Cited on pages 28 and 29.)
- Rafael Poyiadzi, Kacper Sokol, Raul Santos-Rodriguez, Tijl De Bie, and Peter Flach. FACE: Feasible and

- Actionable Counterfactual Explanations. In AAAI/ACM Conference on AI, Ethics, and Society, 2020. (Cited on page 9.)
- Inioluwa Deborah Raji, Morgan Klaus Scheuerman, and Razvan Amironesei. You Can't Sit With Us: Exclusionary Pedagogy in AI Ethics Education. In ACM Conference on Fairness, Accountability, and Transparency, 2021. (Cited on page 67.)
- Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. Why Should I Trust You? Explaining the Predictions of Any Classifier. In ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, 2016a. (Cited on pages 29, 36, 49, and 71.)
- Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. Model-Agnostic Interpretability of Machine Learning. In ICML Workshop on Human Interpretability in Machine Learning, 2016b. (Cited on page 45.)
- Chris Russell. Efficient Search for Diverse Coherent Explanations. In ACM Conference on Fairness, Accountability, and Transparency, 2019. (Cited on pages 9 and 49.)
- Jeffrey Saltz, Michael Skirpan, Casey Fiesler, Micha Gorelick, Tom Yeh, Robert Heckman, Neil Dewar, and Nathan Beard. Integrating Ethics within Machine Learning Courses. ACM Transactions on Computing Education, 2019. (Cited on page 67.)
- Wojciech Samek, Grégoire Montavon, Andrea Vedaldi, Lars Kai Hansen, and Klaus-Robert Müller. Explainable AI: Interpreting, Explaining and Visualizing Deep Learning. Springer, 2019. (Cited on page 28.)
- Michael Sejr Schlichtkrull, Nicola De Cao, and Ivan Titov. Interpreting Graph Neural Networks for NLP With Differentiable Edge Masking. In *International Conference on Learning Representations*, 2021. (Cited on pages 28 and 29.)
- Thomas Schnake, Oliver Eberle, Jonas Lederer, Shinichi Nakajima, Kristof T. Schütt, Klaus-Robert Müller, and Grégoire Montavon. XAI for Graphs: Explaining Graph Neural Network Predictions by Identifying Relevant Walks. *arXiv*, 2020. (Cited on page 30.)
- Lisa Schut, Oscar Key, and Rory McGrath. Generating Interpretable Counterfactual Explanations By Implicit Minimisation of Epistemic and Aleatoric Uncertainties. In *International Conference on Artificial Intelligence and Statistics*, 2021. (Cited on page 28.)
- D. Sculley, Jasper Snoek, Alexander B. Wiltschko, and Ali Rahimi. Winner's Curse? On Pace, Progress, and Empirical Rigor. In *International Conference on Learning Representations*, 2018. (Cited on page 69.)
- Boris Sharchilev, Yury Ustinovsky, Pavel Serdyukov, and Maarten de Rijke. Finding Influential Training Samples for Gradient Boosted Decision Trees. In *International Conference on Machine Learning*, 2018. (Cited on page 48.)
- Hong Shen, Wesley H. Deng, Aditi Chattopadhyay, Zhiwei Steven Wu, Xu Wang, and Haiyi Zhu. Value Cards: An Educational Toolkit for Teaching Social Impacts of Machine Learning through Deliberation. In ACM Conference on Fairness, Accountability, and Transparency, 2021. (Cited on pages 68 and 78.)
- Natasha Singer. Tech's Ethical 'Dark Side': Harvard, Stanford and Others Want to Address It. The New York Times: https://www.nytimes.com/2018/02/12/business/computer-science-ethics-courses.html, 2018. (Cited on page 67.)
- Michael Skirpan, Nathan Beard, Srinjita Bhaduri, Casey Fiesler, and Tom Yeh. Ethics Education in Context: A Case Study of Novel Ethics Activities for the CS Classroom. In *ACM Technical Symposium on Computer Science Education*, 2018. (Cited on pages 67, 68, and 78.)
- Suraj Srinivas, Akshayvarun Subramanya, and R. Venkatesh Babu. Training Sparse Neural Networks. In CVPR Workshops, 2017. (Cited on pages 31 and 32.)
- Ilia Stepin, Jose M Alonso, Alejandro Catala, and Martín Pereira-Fariña. A Survey of Contrastive and Counterfactual Explanation Generation Methods for Explainable Artificial Intelligence. *IEEE Access*, 2021. (Cited on pages 28 and 30.)
- Robert Stojnic. Papers with code partners with arxiv. Medium: https://medium.com/paperswithcode/papers-with-code-partners-with-arxiv-ecc362883167, 2020. (Cited on pages 69 and 73.)
- Jonathan M. Stokes, Kevin Yang, Kyle Swanson, Wengong Jin, Andres Cubillos-Ruiz, Nina M. Donghia, Craig R. MacNair, Shawn French, Lindsey A. Carfrae, Zohar Bloom-Ackermann, Victoria M. Tran, Anush Chiappino-Pepe, Ahmed H. Badran, Ian W. Andrews, Emma J. Chory, George M. Church, Eric D. Brown, Tommi S. Jaakkola, Regina Barzilay, and James J. Collins. A Deep Learning Approach to Antibiotic Discovery. Cell, 2020. (Cited on page 28.)
- Jiawei Su, Danilo Vasconcellos Vargas, and Sakurai Kouichi. One Pixel Attack for Fooling Deep Neural Networks. *IEEE Transactions on Evolutionary Computation*, 2019. (Cited on page 9.)
- Lichao Sun, Yingtong Dou, Carl Yang, Ji Wang, Philip S. Yu, Lifang He, and Bo Li. Adversarial Attack and Defense on Graph Data: A Survey. *arXiv*, 2018. (Cited on page 30.)
- Thomas Douglas Victor Swinscow. Statistics at Square One. BMJ Publishing Group, 1997. (Cited on pages 51

- and 52.)
- Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus. Intriguing Properties of Neural Networks. In *International Conference on Learning Representations*, 2014. (Cited on page 9.)
- Maartje ter Hoeve, Mathieu Heruer, Daan Odijk, Anne Schuth, Martijn Spitters, and Maarten de Rijke. Do News Consumers Want Explanations for Personalized News Rankings? In *FATREC Workshop on Responsible Recommendation*, 2017. (Cited on page 56.)
- Gabriele Tolomei, Fabrizio Silvestri, Andrew Haines, and Mounia Lalmas. Interpretable Predictions of Tree-based Ensembles via Actionable Feature Tweaking. In *ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2017. (Cited on pages 8, 9, 10, 14, 16, 30, and 48.)
- John W. Tukey. Exploratory Data Analysis. Addison-Wesley, 1977. (Cited on page 50.)
- UCI. Wine Quality Data Set. https://archive.ics.uci.edu/ml/datasets/Wine+Quality, 2009. (Cited on page 14.)
- UCI. Online Shoppers Intention Dataset. https://archive.ics.uci.edu/ml/datasets/Online+Shoppers+Purchasing+Intention+Dataset, 2019. (Cited on pages 14 and 15.)
- Berk Ustun, Alexander Spangher, and Yang Liu. Actionable Recourse in Linear Classification. In ACM Conference on Fairness, Accountability, and Transparency, 2019. (Cited on pages 9, 28, 30, and 40.)
- Arnaud Van Looveren and Janis Klaise. Interpretable Counterfactual Explanations Guided by Prototypes. In Nuria Oliver, Fernando Pérez-Cruz, Stefan Kramer, Jesse Read, and Jose A. Lozano, editors, European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases, 2021. (Cited on page 9.)
- Sahil Verma, John Dickerson, and Keegan Hines. Counterfactual Explanations for Machine Learning: A Review. arXiv, 2020. (Cited on pages 28 and 30.)
- Minh N. Vu and My T. Thai. PGM-Explainer: Probabilistic Graphical Model Explanations for Graph Neural Networks. In *Advances in Neural Information Processing Systems*, 2020. (Cited on pages 28 and 29.)
- Sandra Wachter, Brent Mittelstadt, and Chris Russell. Counterfactual Explanations Without Opening the Black Box: Automated Decisions and the GDPR. *Harvard Journal of Law and Technology*, 2018. (Cited on pages 8, 9, 11, 30, 31, and 49.)
- Oliver Wieder, Stefan Kohlbacher, Mélaine Kuenemann, Arthur Garon, Pierre Ducrot, Thomas Seidel, and Thierry Langer. A Compact Review of Molecular Property Prediction with Graph Neural Networks. Drug Discovery Today: Technologies, 2020. (Cited on page 28.)
- Yutong Xie, Chence Shi, Hao Zhou, Yuwei Yang, Weinan Zhang, Yong Yu, and Lei Li. MARS: Markov Molecular Sampling for Multi-objective Drug Discovery. In *International Conference on Learning Representations*, 2021. (Cited on page 28.)
- Yongxin Yang, Irene Garcia Morillo, and Timothy M. Hospedales. Deep Neural Decision Trees. In *ICML Workshop on Human Interpretability in Machine Learning*, 2018. (Cited on page 10.)
- Burak Yildiz, Hayley Hung, Jesse H Krijthe, Cynthia CS Liem, Marco Loog, Gosia Migut, Frans A Oliehoek, Annibale Panichella, Przemysław Pawełczak, Stjepan Picek, et al. ReproducedPapers.org: Openly Teaching and Structuring Machine Learning Reproducibility. In *International Workshop on Reproducible Research in Pattern Recognition*, 2021. (Cited on page 68.)
- Rex Ying, Dylan Bourgeois, Jiaxuan You, Marinka Zitnik, and Jure Leskovec. GNNExplainer: Generating Explanations for Graph Neural Networks. In *Advances in Neural Information Processing Systems*, 2019. (Cited on pages 28, 29, 34, 35, 36, and 37.)
- Hao Yuan, Jiliang Tang, Xia Hu, and Shuiwang Ji. XGNN: Towards Model-Level Explanations of Graph Neural Networks. *arXiv*, 2020a. (Cited on page 29.)
- Hao Yuan, Haiyang Yu, Shurui Gui, and Shuiwang Ji. Explainability in Graph Neural Networks: A Taxonomic Survey. *arXiv*, 2020b. (Cited on pages 28, 29, and 36.)
- Hao Yuan, Haiyang Yu, Jie Wang, Kang Li, and Shuiwang Ji. On Explainability of Graph Neural Networks via Subgraph Explorations. In *International Conference on Machine Learning*, 2021. (Cited on pages 28 and 29.)
- Marinka Zitnik, Monica Agrawal, and Jure Leskovec. Modeling Polypharmacy Side Effects with Graph Convolutional Networks. *Bioinformatics*, 2018. (Cited on page 28.)

Try to combine abstracts from each chapter

FOCUS abstract: Model interpretability has become an important problem in machine learning (ML) due to the increased effect that algorithmic decisions have on humans. Counterfactual explanations can help users understand not only why ML models make certain decisions, but also how these decisions can be changed. We frame the problem of finding counterfactual explanations as a gradient-based optimization task and extend previous work that could only be applied to differentiable models. In order to accommodate non-differentiable models such as tree ensembles, we use probabilistic model approximations in the optimization framework. We introduce an approximation technique that is effective for finding counterfactual explanations for predictions of the original model and show that our counterfactual examples are significantly closer to the original instances than those produced by other methods specifically designed for tree ensembles.

CF-GNN abstract: Given the increasing promise of graph neural networks (GNNs) in real-world applications, several methods have been developed for explaining their predictions. Existing methods for interpreting predictions from GNNs have primarily focused on generating subgraphs that are especially relevant for a particular prediction. However, such methods are not counterfactual (CF) in nature: given a prediction, we want to understand how the prediction can be changed in order to achieve an alternative outcome. In this work, we propose a method for generating CF explanations for GNNs: the minimal perturbation to the input (graph) data such that the prediction changes. Using only edge deletions, we find that our method can generate CF explanations for the majority of instances across three widely used datasets for GNN explanations, while removing less than 3 edges on average, with at least 94% accuracy. This indicates that our method primarily removes edges that are crucial for the original predictions, resulting in minimal CF explanations.

MC-BRP abstract: In various business settings, there is an interest in using more complex machine learning techniques for sales forecasting. It is difficult to convince analysts, along with their superiors, to adopt these techniques since the models are considered to be "black boxes," even if they perform better than current models in use. We examine the impact of contrastive explanations about large errors on users' attitudes towards a "black-box" model. We propose an algorithm, Monte Carlo Bounds for Reasonable Predictions. Given a large error, MC-BRP determines (1) feature values that would result in a reasonable prediction, and (2) general trends between each feature and the target, both based on Monte Carlo simulations. We evaluate on a real dataset with real users by conducting a user study with 75 participants to determine if explanations generated by MC-BRP help users understand why a prediction results in a large error, and if this promotes trust in an automatically-learned model. Our study shows that users are able to answer objective questions about the model's predictions with overall 81.1% accuracy when provided with these contrastive explanations. We show that users who saw MC-BRP explanations understand why the model makes large errors in predictions significantly more than users in the control group. We also conduct an in-depth analysis of the difference in attitudes between Practitioners and Researchers, and confirm that our results hold when conditioning on the users' background.

EAAI abstract: In this work, we explain the setup for a technical, graduate-level course on Fairness, Accountability, Confidentiality, and Transparency in Artificial Intelligence (FACT-AI) at the University of Amsterdam, which teaches FACT-AI concepts through the lens of reproducibility. The focal point of the course is a group project based on reproducing existing FACT-AI algorithms from top AI conferences and writing a corresponding report. In the first iteration of the course, we created an open source repository with the code implementations from the group projects. In the second iteration, we encouraged students to submit their group projects to the Machine Learning Reproducibility Challenge, resulting in 9 reports from our course being accepted for publication in the ReScience journal. We reflect on our experience teaching the course over two years, where one year coincided with a global pandemic, and propose guidelines for teaching FACT-AI through reproducibility in graduate-level AI study programs. We hope this can be a useful resource for instructors who want to set up similar courses in the future.

Sažetak

Ovo je sažetak ...

Samenvatting

Dit is de samenvatting ...